

The EE3 Smart House

Exploring the application of machine learning in the context of home automation

Grubb, Sebastian
sg3510@ic.ac.uk

Murray, Michael
mm510@ic.ac.uk

Janssen, Loek
lfj10@ic.ac.uk

Zhao, Tony
tz1010@ic.ac.uk

Mason, Jonathan
jhm210@ic.ac.uk

Singleton, Matthew
ms8210@ic.ac.uk

van Beek, Jonathan
jv510@ic.ac.uk

Hepburn, Connel
ch3810@ic.ac.uk

June 22, 2013

Department of Electrical and Electronic Engineering
Imperial College London
London SW7 2AZ

Executive Summary

With modern societies continued demand for energy, dwindling natural resources and rising prices, technologies that can curb wasteful expenditure are becoming ever more relevant. In 2009, 28% of the total energy produced in the UK was consumed by households, 65.7% of which was derived from heating alone. Therefore there is a need for a product which can improve the efficiency of household central heating systems, and as a by product save people money. Our solution is an automated system which not only adjusts the home environment to suit the needs of the user, but also minimises wasteful energy consumption. With a unique machine learning approach, the system we propose in this report aims to revolutionise the way in which home environments are controlled.

The scope of existing home automation systems is varied, some providing only remote control while others just data processing. Autohome is a fully integrated system using a unique machine learning approach, that not only automates thermostat settings, but also provides user interaction and remote appliance control. As a result Autohome offers a more comprehensive service than any of its competitors, presenting its users with greater functionality and potential energy savings.

The Autohome solution can be broken down into two key layers; hardware and software. The hardware layer consists of a network of sensors situated in the users home that collects data on the internal conditions. The software layer sits on a server and processes the collected data to infer thermostat predictions. User interaction with the system is achieved through a mobile app and website. This allows the user to manually adjust the thermostat, remotely control appliances and view feedback statistics on their financial and energy expenditure. As a byproduct of making the user more aware of their energy usage, Autohome should encourage users to become more energy conscious.

Home automation is an emerging market which is expected to grow 16.1% year-on-year worldwide. However, for example, currently only 0.1% of UK and 5% of US homes have a home automation system. Therefore there is huge untapped potential in this market, with as yet no cemented market leader. Considering this, and also combining modern society's concerns over fossil fuel emissions with the current trend in smart interconnected devices, the conditions are ideal for Autohome to enter the market.

Contents

1	Introduction: <i>Problem Identification and Market Feasibility</i>	4
1.1	Problem Identification	4
1.2	Market and Product Feasibility	4
1.2.1	Product Feasibility	4
1.2.2	Market Feasibility	5
1.2.3	Competitors	5
2	Top Level Design	6
3	Modular Design	7
3.1	Hardware	7
3.1.1	mbed	7
3.1.2	Wireless Communications	7
3.1.3	Sensors	9
3.1.4	Other Hardware	10
3.2	Software	11
3.2.1	Server Architecture and Technologies	12
3.2.2	Hidden Markov Model for Occupation State Detection	13
3.2.3	Timetable Prediction	13
3.2.4	Machine Learning for Thermostat Setting Prediction	14
3.3	User Interface	18
3.3.1	Website	18
3.3.2	Mobile App	18
4	System Analysis and Use	19
4.1	Savings Performance	19
4.2	Hardware Performance	21
4.3	Software Performance	21
4.3.1	Timetable Prediction	21
4.3.2	Website Usage	22
4.4	System Practicality	22
4.5	System Unit Cost	23
5	Further Development	24
5.1	LCD Screen	24
5.2	Modular Network Setup	24
5.3	Smart Door Lock	24
6	Conclusion	25
7	References	26
8	Appendix	28
8.1	Hardware	28
8.1.1	mbed	28
8.1.2	Wireless Power Monitoring	46
8.2	Software	48
8.2.1	Server Code	48
8.2.2	Timetable Prediction	49
8.2.3	Gaussian Process Regression	54
8.3	Costs	69

Section 1

Introduction

Problem Identification and Market Feasibility

1.1 Problem Identification

Modern society's continued and growing need for energy, combined with dwindling fossil fuel reserves and environmental concerns around CO₂ emissions, has created an environment in which technology that can limit the wasteful expenditure of energy can thrive. In 2009 UK households consumed 501TWh accounting for 28-32% of the UK's total energy expenditure [1]. Central heating undeniably plays its role here, and following the current trend of a 3.7 degree centigrade rise in the internal temperature of houses since 1970[2], and rising numbers of UK households its contribution is only set to increase. As a result there is a need for products that can improve the efficiency of a household heating system.

Current household heating systems are controlled manually via a thermostat, which is neither convenient for the user or sensitive to the actual environmental conditions. In today's busy world few people have the time or discipline to consistently alter their thermostat to fit their own dynamic schedule or the variable nature of the weather. This is not only wasteful in terms of energy but is also very expensive. Therefore there is a niche for a system that can learn and adapt to the chaotic nature of people's lives and the weather.

The proposed solution to this problem is to create a machine learning system that once trained by the user, intelligently checks whether the occupants of the house are in, out or asleep; and sets the thermostat setting in accordance with this and the external conditions. This system should therefore save energy and make the user experience more convenient while not compromising on the comfort of the user. As a by product of saving energy the system should also save the user money in the long term; an important consideration is that in the UK 2.6 million households are in fuel poverty, and hence spend at least 10% of their income on heating [3].

1.2 Market and Product Feasibility

Objects are increasingly gaining the ability to communicate and sense their environment, forming what is broadly defined as the Internet of Things. This ability to collect and transmit a vast amount of data in real-time gives rise to new insights, allowing for better day-to-day decisions to be made.

The home automation market forms a subset of the Internet of Things market, with both expected to grow worldwide in 2013 with a CAGR of 16.1%[4] and CAGR¹ of 33.2%[5], respectively. Thus it is the ideal time to build a home automation product when so many enablers exist.

1.2.1 Product Feasibility

The first enabler of a home automation product is the reduction in cost of low power processing components. For example the processor used for AutoHome costs £6.39 per unit[6], with sensors often costing less than £1 which means that when assembled an affordable sub £150 build price is expected. The second enabler is the increased penetration of smartphones and wireless technology that allows the possibility of an elegant home automation implementation, along with increased consumer demand for such technological solutions. Thirdly is the rise of cloud computing which is the main enabler of AutoHome, as data collected in real-time cannot be as efficiently processed on the mbed as the dataset becomes larger (sensor data being sampled at 5 minute intervals leads to 8640 data-points per house per month) - this is especially true of the computationally intensive machine learning algorithms. Computations are instead done on remote servers (such as Amazon's EC2 compute solution used for AutoHome), which are often more powerful and manage computing loads more efficiently. Additionally current

¹Compounded Annual Growth rate

server solutions offer competitive rates priced at \$0.060 per hour and above, relatively affordable operating costs for computing power.

1.2.2 Market Feasibility

With 17.5 million owner occupied homes in the UK and 87.4 million owner occupied homes in the US the home automation market has a huge potential. However only 0.1% and 5% of homes, in the UK and US respectively, have such home automation systems. This poor market penetration is due to multiple factors - such as the current state of home automation products, with many of them only offering remote automation and not smart automation. The key difference being that remote automation requires constant human interaction to control the house's settings (though scheduling is possible it is not truly smart), whereas smart automation provides a frictionless experience, constantly adapting to the user's needs. A second reason is the high prices (over-£300) of currently available systems[7], which are also described as clunky and tiresome to install. Thirdly is the lack of incentives, as most systems merely act as an additional way to control the house - which amounts to paying a premium to install extra light switches. There is thus little value seen in the cost of such devices, leaving room for a category of affordable, frictionless and smart home automation products which give the consumer clear value - such as saving money and energy.

Additionally recent UK policy has mandated that all UK energy providers must retrofit a smart meter in homes by 2020[8], meaning there is an additional market involving teaming up with energy providers to provide such systems.

The main market segments are homeowners, where retrofitting this device will be necessary. The other major segment is that of new homes, where there is a focus on delivering modern and energy conscious habitats - meaning that a deeper integration is possible, leading to even more efficient system.

1.2.3 Competitors

Many companies exist in the home automation market, such as AMX, Control4, iControl Networks and Vivint. However most of these only offer remote automation, as explained earlier, though they cover more components of a house, such as window shutter, TV control, house music and personal CCTV. The table below gives a breakdown.

	AutoHome	Control4	Vivint	Nest Thermo-stat	Opower
Cost	£150-£200	<£650	Starts at £38/month	£161	N/A (B2B company)
Features	-Smart thermostat -Advanced house state detection -lighting and smart meter offered -focused on energy savings -easily installable	-Availability by modules (starter kit \$1000) -Work across residential, commercial and hospitality industries -Offer in depth control including security solutions -Requires contractor for installation	-Offer Home Automation, Energy Management and Home Security solutions -Provide intuitive solutions -No real "smart" automation -Requires contractor for installation	-Smart temperature control -Strong brand in part due to being made by ex-Apple engineers -Focused on saving money -Does not require contractor for installation	-Very good at energy analytics -Sold via energy companies -Provides intelligent software for thermostats -Purely focused on energy
Comparison	N/A	AutoHome smarter, focusing more on energy saving features	Smarter than Control4 but not to the extent AutoHome offers, no monthly fee for AutoHome.	Possibly main competitor. Offers similar "intelligent" algorithms but does not use as much data to regress on as AutoHome.	Competes on the algorithm/machine learning side but offers no other home automation solutions.

Section 2

Top Level Design

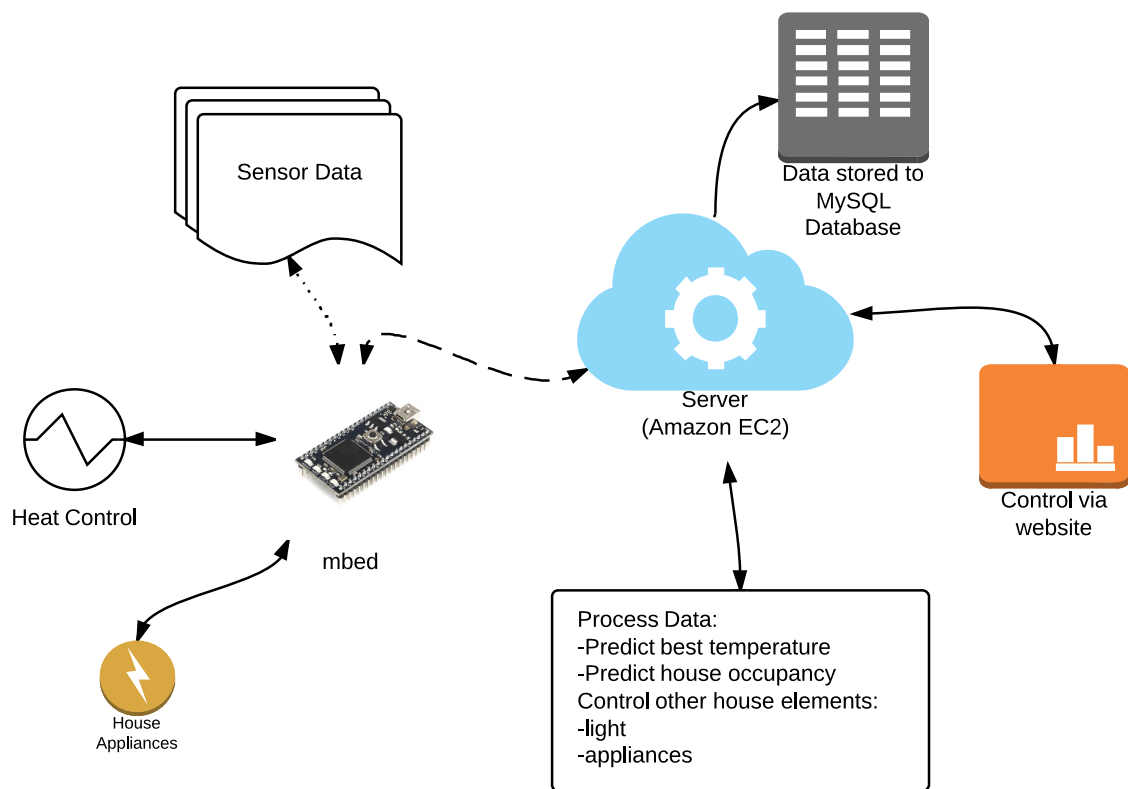


Figure 2.1: Overall System Design

Section 3

Modular Design

3.1 Hardware

For the first stage a prototype would be developed, which would include temperature and humidity measurement; an xbee module and a control interface; as well as the ability for the microprocessor to connect to the internet. To facilitate prototyping an 'mbed', prototyping board was used. This board is similar to a microcontroller breakout board but with many extra features, including a USB to FTDI connection for ease of programming, reset buttons and LEDs. This allowed for quick development and allowed basic prototyping to go ahead at an early stage. Each module was then developed in isolation using the mbed board, and then the final programme was written to combine the various elements together, after which the final link up of software and hardware was devised.

3.1.1 mbed

The 'mbed' prototyping board contains an LPC1768 processor which is a 32-bit ARM cortex-M3 microcontroller and the board can be powered either through a USB connection or a 4.5-9.0V input voltage. Figure 3.1 shows the various pins that are available with the mbed board, including a regulated 3.3V supply voltage which is used to power several of the main components.

3.1.2 Wireless Communications

As discussed in the top-level design, two wireless protocols would play pivotal roles in bringing the entire system together, with the standard wireless local area network (or Wi-Fi) to communicate with the internet via wireless routers which are now found within over 70% of UK households [25]. Then the xbee protocol was used to communicate with the various sensors and components around the house back to the microprocessor, due to both its low-power usage and the ability to easily create mesh networks. The mbed could then be used to pass the information and commands between the two networks (and hence our servers) as required.

Wi-Fi

As stated above there is a reasonable expectation that any home buying our smart thermostat will have wireless internet, both due to the target market and the high proportion of households with wireless routers. The roving networks Wifly RN-XV module was chosen due to its easy integration into the mbed system and the ability to easily call an ultra-low sleep mode during periods of decreased activity. One slight disadvantage was the inability of the module to join enterprise security networks at the current moment of time, with an updated firmware expected to make such connections possible within a year.

The Wifly module was then set up as shown in the overall circuit diagram (Appendix Section 8.1.1), using one of the 3 serial pair (RX and TX) connections available on the mbed. Simple ASCII commands using a serial console (in this case PuTTY) were used to experiment with the Wifly module, allowing demonstrable access to the internet which was tested using several standard ISP provided home routers.

Communicating with the server could be done via several methods, such as http and the new websockets. Both protocols were experimented, with websockets ultimately being chosen, this was due to the ability to have constant 2-way communication between the mbed and server, which could be set-up simply and easily. Using pre-existing libraries the correct SSID and WPA2-PSK code was placed in the Wifly, and then at the beginning of the programme the module is commanded to join the network (see appendix for code). One disadvantage that was




Figure 3.1: mbed NXP LPC1768 pinout diagram

found was that if the module failed to join a network over several resets (of the mbed) it would crash and have to be rebooted, either by power cycling or using command mode via PuTTY.

XBee

The XBee modules play a paramount role in servicing communications between all the individual sensors and the MBED which is at the heart of everything. We chose the XBee as we believed that it provided the best features for our design. The XBee ZNET protocol can be used to create a mesh network from all of the XBees. This is ideal for application in a household environment where consumers could add extra sensors which could possibly be placed a significant distance away from the central unit.

Figure 3.2 helps show how the mesh network operates. Our design incorporates one XBee co-coordinator on the central unit with all of the peripherals being configured as router XBees. The XBees which we used had the specifications listed below; the model which was used is the standard XBee not the XBee PRO.

The XBee specifications allow a 40m indoor range which is more than enough to cover most households. On top of this is the fact that each XBee can use its 40m range to form a mesh network allowing us to greatly increase the effective range of the network. The transmission power is far less than the WiFly equivalent. The power of the XBee is only 2mW even in boost mode vs. an average transmission power of 446mW of the lowest output power rating on the WiFly module in 802.11g mode.

There are two modes of communication to the XBee module, one is called transparent mode or AT mode and the other is the API (Application Programming Interface) mode. Our design is based upon the API mode as it allows greater flexibility and is significantly faster. This could be seen that when issuing commands to the XBee via AT, such as a change to a digital output pin level, there would be a noticeable amount of lag until transition occurred which was not found with API firmware. One other option which the API firmware provided the ability to easily address individual XBees in order to send them commands. This is critical as unique commands are associated with each peripheral type. The AT firmware requires that two registers be set with separate commands to change the destination address when sending to a specific XBee. This creates additional processing load for both the XBee and the central processing unit as well as reducing the speed at which commands can be sent to the XBees. API mode allows 64 bit addressing of XBees within an API frame, allowing a more efficient and convenient method of sending commands with the API framework.

The XBee uses a custom API framework which we created for the purposes of this project. This allows the MBED to construct basic API frames and send them out over its UART port. We have currently only implemented the frames which are required for our project to work though as it is a time consuming process to write and test the operation of a specific command. Currently the framework supports: Requests for the PAN ID, Changing the PAN ID, Sending data packets to other XBees, Requesting the value of all analogue and digital input pins, setting digital levels on the digital output pins.

Details of the framework, its use with different modules, the custom developed API and observed issues can be found in the appendix section 8.1.1.

Limitations and improvements

Had there been more time to debug and implement a more robust framework it would have been possible to improve it in a few ways. The first is to get a fully working UART interrupt such that everything can then be called on demand and not consume power when it's not needed as well as providing more control on the data coming into the servers. This could be useful in order to throttle some of the connections if the server architecture wouldn't be able to handle the load. One more interesting thing would be to write custom firmware for the XBees. By doing this we could offload menial tasks from the MBED and put them onto the XBee. This could be useful if we wanted to convert the power calculations from their raw values to their actual values. This could then be offloaded from the MBED. This would help as the MBED uses a lot more power in its current form than one of the XBee modules. The power drawn has been measured at around 0.7W under normal operating conditions which is considerably higher than the XBee. By taking the load off the mbed we can activate and optimise a sleep routine for the MBED.

This can also be further advanced by changing the MBED with another model such as the one based on the Cortex M0+ from Freescale. This would allow even more power efficient operation as long as there was still enough processing power to run everything. The issue with the M0 is that it does have less connectivity options such as UARTs meaning that some external hardware might have to be used to compensate, thereby increasing the cost.

Lastly the MBED has its own RTOS¹. By using an RTOS it would be possible to thread the web socket operations and the xbee operations in order to get more throughput from the microcontroller as well as massively reducing latency as well. This would come in useful with the power routine and its large wait cycle, allowing the

Legend:
C - Coordinator
R - Router
E - End device




Figure 3.2: Diagram depicting the arrangement of the mesh network.

¹Real Time Operating System

websocket to be serviced in the meantime. This would however increase complexity and thus would have taken more time to implement and debug.

3.1.3 Sensors

Temperature sensor

In order to measure temperature an 8-pin digital sensor was used that has a 2 wire serial interface with the microcontroller, using the I2C pins p28 and p27 upon the mbed board. This was the TMP75AIDR from Texas Instruments and for sensing the temperature, the chip itself is used with thermal paths running through the packaging. This is then converted internally to a digital temperature output with an accuracy of $\pm 0.5^\circ\text{C}$, which is then communicated to the microcontroller via I2C.

The I2C is a multimaster bus which allows easy digital communication with multiple devices if required, and hence the slave address must be set using 3 pin inputs which should match up to the chosen address in the software. Setting the 3 pins to ground gives the address 0x48 which we can see in the temperature configuration section of the code. Using libraries previously developed for mbed, a simple programme was devised (see appendix) which would only use a 9-bit resolution (set in the temperature configuration code) as a resolution of 0.0625°C (12-bit resolution) was deemed unnecessary. The chip itself was an SMD (surface mount device) and hence a small breakout was constructed to allow ease of use with breadboards.

Humidity sensor

Within the thermostat algorithm, internal humidity would play a role and so the low voltage HIH-5030 humidity sensor was acquired, also being an SMD model a small breakout board was required once again. The device has only 3 pins, a V_{cc} , ground, and analogue output which when measured can be used to determine the relative humidity (RH) of the atmosphere to $\pm 3\%$.

Using the information supplied in the datasheet two equations were identified, one which could supply the sensor RH from the output voltage, and another to determine the true RH as the sensor RH is also affected by the temperature. Hence the temperature sensor is also required, the final equations within the mbed programme were then

$$SensorRH = \left(\frac{Voltage \times 5}{3.3} - 0.1515 \right) / 0.00636 \quad (3.1)$$

$$TrueRH = \frac{SensorRH}{1.0546 - (0.00216 \times Temperature)} \quad (3.2)$$

The final true relative humidity would then be calculated within the mbed, which would use an analogue-to-digital pin (ADC) to measure the output voltage of the sensor, the result can then be displayed and returned to the servers as required.

Wireless Power Monitoring

The ability to monitor the power consumption of a house in real time is useful in both making the user aware of the impact their habits have on their electricity bill and also to justify any saving the auto-home enables. This is by no means a new idea, and the British government, among many others, intend on a smart meter roll out scheme whereby every home in the UK will have a real time power monitor, or "smart meter", by 2020[9]. With this same power information, it is also possible to detect the appliances that are running in the house with non-intrusive load monitoring[10] (NILM), which can indicated unnecessary power usage when nobody is in the house, and switch these off to save power. With this considered, a power monitor is a useful part of the system, and was designed to make accurate measurements with minimum power usage and transmit this to the central unit.

Although measuring power seems a trivial exercise, at AC there are many ways of doing it. One way of making the measurement is calculating the power by summing many instantaneous voltage and current products over a period and divide by the number of samples. Although this is probably the most reliable and accurate method, it is also rather computationally intensive for a low powered 8-bit micro-controller. With the assumption that the voltage and current will have a reasonably consistent shape (sinusoidal), another simpler method is to measure the peak-peak voltage and current, and translate to RMS with a constant factor. However, this is quite a large assumption to make, and for many appliances such as switch mode laptop supplies, the current signal will be very different to that of a light bulb, but this will be less significant when looking at the total draw of a house giving an overestimation of the power usage, which isn't a bad thing.

With the peak-peak voltage/current measurements and a constant factor (for true sinusoids $\frac{\sqrt{2}}{2}$), the conversions can be made to RMS and the apparent power calculated from the product of the two. However, to be able to accurately calculate the dissipative real power, or characterise a load in terms of individual appliance signatures, the phase information must also be recorded. Again, for simplicity and ease of debugging, a comparator \rightarrow XOR \rightarrow low-pass filter phase difference detector was used to give a third ADC input corresponding to the angle between

the voltage and current. Again, by only a simple constant conversion factor, the angle and power factor can easily be calculated.

Motion Sensor

An important feature of the home automation project is the ability to detect whether there is anyone within the house. To do this we are using multiple passive infrared motion sensors, located within key locations, to monitor movement. This will provide one of the main sources of data for our learning algorithms. Our sensor itself is a long range infrared sensor with a detection range of up to 12m. With a detection angle of up to 102 degrees, this ensures that for the average sized house, all motion within a room will be detected (if placed in one corner). After an initial stabilisation time of roughly 30 seconds, the sensor will output a digital high when the detecting target is present and a digital low otherwise. This is fed straight to the XBee and then to the mbed for processing.

The method of detection for the sensor is via infrared radiation. This is perfectly suited for human detection from the radiation given out by their body temperature. It also rules out the potential of false triggering from wind blowing on open doors or curtains. However, this does not remove the problem of detection from pets. This is dealt with by the algorithm instead. The sensor consists of multiple polarized detection zones, each zone surrounded by 4 of opposite polarity. When the target moves across these detection zones, the combined polarity varies, hence movement is detected. However, a problem may occur if a target enters a positive and negative detection zone at the same time. In this case the signals will cancel each other and no detection notified. This will usually only occur near the maximum detection range, where the target is smaller in respect to the detection zone. Therefore it is not a significant problem as most rooms would not reach this distance.

Some other features of the sensor include a metallic can which encloses the sensing circuit. This increases the signal to noise ratio of the sensor. The result provides protection from the false detection of external electromagnetic fields caused by mobile phones and other electronic devices. The simplicity of the device along with the negligible power drawn when no detection is present means that very little total power is consumed during operation. The




Figure 3.3: Space detection

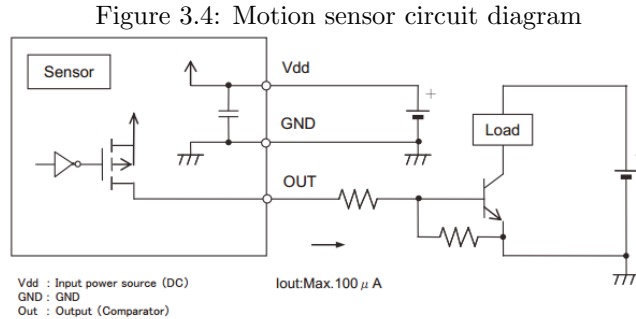


Figure 3.4: Motion sensor circuit diagram

circuit is very simple and can be seen in figure 3.4. The load is replaced by the XBee and resistor values set to output current/voltage. The input power source is the standard 3.3V used for our XBee and mbed.

3.1.4 Other Hardware

Lighting

As part of the focus to improve user comfort and energy savings, another important feature of the home automation system is the capability to remotely control lights around the house. A large portion of a homeowner's electricity bill is contributed by lighting in the house. In the U.S, this figure is roughly 11 percent of a household's entire energy budget, and is significantly greater for commercial buildings [12]. Therefore only keeping the lighting on when necessary can provide substantial savings.

Our main feature is to keep track of all lighting in the house, allowing the user to see the states of the lights on either an app on their phone or via the website. The user can then change the state of the lights to on, off or auto accordingly. Notifications can also be set to alert the user if any light has been on for a certain length of time. The auto setting will access the motion sensors in the selected room to set those lights to be motion detected. People tend to have various preferences and opinions when it comes to motion detecting lights; therefore our system allows the user full control to choose what suits them best in each





Figure 3.5: Triac diode diagram

room. For example, when in the living room watching TV, the auto setting may not be appropriate as the light will alternate between on/off due to the lack of movement. But the setting can be very convenient for a toilet, as it is used for shorter periods. All these features can ensure no lights are left switched on by accident.

Figure 3.6: Triac V/I characteristics



To remotely control the light switch, we use our mbed to send a signal via the zigbee wireless to an XBee module next to the switch. The XBee will then output a voltage as an input to a XOR gate, with the other being the manual switch located on the wall. Toggling either input will trigger an optoisolator, which will control the high power triac switch, turning it on/off. The triac is connected directly to the mains line of the light bulb, and will conduct only when triggered. The state of the light bulb will then be sent back via the XBee, updating the website and app.

The characteristics of a triac are ideal for our purpose of switching on/off the mains voltage. As shown in figure 3.5, it is made on two diodes with a gate connection leading from the side of MT1. The function can be comparable to a MOSFET, conducting in either direction when the gate requirements are met. However, unlike the MOSFET, the triac can conduct both AC and DC currents. This is vital as the mains voltage is always an AC. The gate of a triac is triggered by a current, rather than voltage for a MOSFET, and will remain conducting until the current drops below the latching current. For an alternating current, this means the triac will switch off twice every cycle. The V/I characteristics can be seen in figure 3.6.

To control the gate current of the triac, a zero-crossing optoisolator is used. The zero-crossing section ensures the gate of the triac is switched on when the current reaches 0. This is important as the phase of the current will be zero, so no noise is generated. Without it, switching could occur at peak voltages, which not only would produce EM noise, but also result in a large $\frac{dV}{dt}$ and potentially blowing the load. The optoisolator also keeps the mains voltage electrically isolated from the XBee circuit. This prevents the mains voltage blowing up the XBee in the case of a short circuit or if the triac blows. Isolation is achieved by simply using an LED to trigger a phototransistor. Hence, the input of the Xbee is attached to pins 1 and 2, and the triac to pins 4 and 6 (figure 3.7).

Not only does remotely controlled lighting reduce energy costs, it can also serve as a tool to provide home convenience for users. With a touch of the phone, no-longer do users need to feel around in the dark for the switch to use the toilet at night, or to get out of bed to switch lights off. Because the feature simply controls the mains supply, it can be easily applied to all electronic appliances. TVs, dishwashers and washing machines are just a few examples of common household appliances that are often left on standby. Switching these off when not in use can provide further cost cuts in the long run.




Figure 3.7: Triac module diagram

3.2 Software

Note: Software was developed in parallel of the hardware, thus no past data collection was available. Instead data made publicly available for the purpose of smart homes research used. The main sources used were the University of Massachusetts Amherst SMART* dataset [11] as well as Washington State’s Tulum dataset [13].

3.2.1 Server Architecture and Technologies

The main server architecture works on a LAMP² server as this provides the system with a database (MySQL) to store all the collected sensor data, a means of processing this data (Python with the numerical NumPy and SciPy libraries) and easily controlling the system (HTTP server with PHP allowing delivery of interactive web-pages).

To enable the website to be able to access the data that is stored in the database, such as the predicted timetable or the past weather information, the server-side programming language of PHP was used. This is the language that is generally used for database access on websites, as it has good support for SQL databases and is easily used in conjunction with HTML. All PHP code is also executed on the server, before being returned to the end-user. This makes it useful as the code must include the database username and password, which would otherwise be passed to the user in the source code, giving anyone access to the database and the ability to alter information. PHP can also be used to create session variables on the server, which allows the creation of a basic login system. Again, as the code is executed server side, a redirect to the login page can be made if the user is not logged in, which will be executed before any of the html is loaded.

The only disadvantage of PHP being executed server side is that it cannot be used to write executable functions on the same page, which is needed when making changes to the timetable and a few other parts of the website. However using ajax from the jquery library, we can call other web pages without leaving the current page, and even pass variables to them which provides a solution for us.

PHP is also used for calling the weather API, to get the current weather conditions. The code gets the contents from a specially generated URL that contains the current weather conditions in JSON format. This can then be decoded, making each weather element easily accessible and these can then be stored into a database. This code is then be called every five minutes using a simple python script running constantly on the server.

Server Software

To control the system a constantly running program, monitoring, processing and relaying all the messages passed is needed. Thus a server-side program which could both send and receive messages to and from system modules was made.

To pass on messages the websocket technology was used. While a relatively new technology, websockets were chosen for their ideal use in real-time applications. This avoids using old 'hacks' such as using HTTP POST and GET requests³ at irregular intervals to achieve a near-realtime connection. This is because websockets provide a full duplex connection with no need of polling⁴ from the client. This would allow the message to turn off a light, or turn the temperature up, to be sent and received nearly instantly - thus being as useful if not more than a real switch.

The main server was written in Python, chosen for its ease of use and good integration with web technologies, using the [tornado](#) framework as code base. The server works by receiving messages in JSON⁵ and passing them on to the appropriate program. For example if a user, via the website page, requests the current house temperature the server receives it and passes it on to the mbed module. Alternatively if the mbed sends the current sensor data to the server a separate script is called to add the values to a database and process this data to predict the most ideal temperature, for example. All the server code can be found in the appendix section 8.2.1.

²The LAMP acronym refers to the first letters of Linux, Apache(a HTTP server), MySQL (database solution), and PHP or Python, the main components to build most web servers.

³HTTP POST and GET requests are the traditional way in which webpages get accessed

⁴i.e. calling the server to ask for a reply

⁵JSON stands for JavaScript Object Notation and is a human readable standard allowing simple data structures to be represented. For example to encode the variable name='Alex' and age=26 the following message would be sent:{"name":"Alex","age":26}

Figure 3.8: Flow diagram of the server's general function

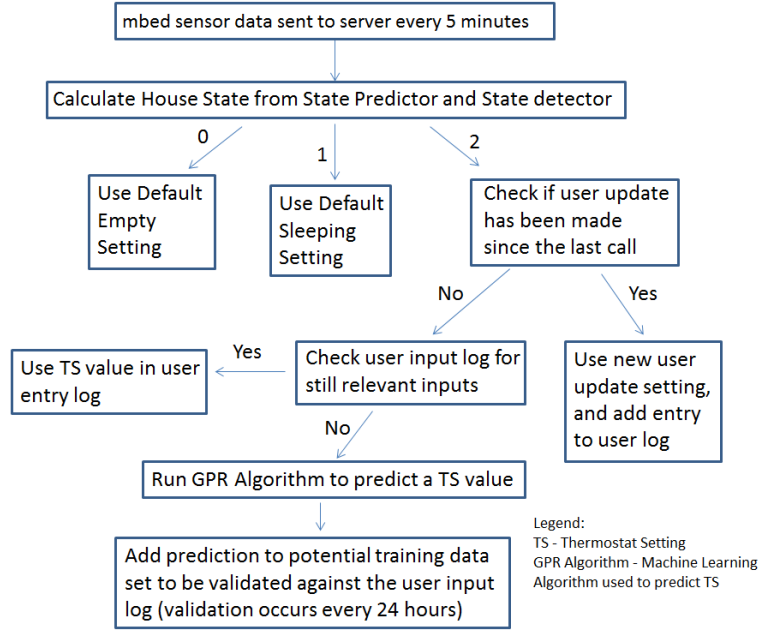


Figure 3.8 is an overview of how the server works when receiving the mbed data to then control the house's thermostat setting.

3.2.2 Hidden Markov Model for Occupation State Detection

3.2.3 Timetable Prediction

While a Hidden Markov Model provides a good way to determine the current state of the house another essential feature to ensure comfort and optimize energy usage is timetable prediction. Different methods were considered such as using the Hidden Markov Model to predict future states, taking an average state of every day of the week or making an ARMA/ARIMA model. The main criteria for choosing a method was an efficient one which could predict to a good degree of accuracy what will happen in about 6 hours in the future. Thus a method suggested by Microsoft Research [14] was implemented and adapted.

The idea behind the proposed solution is to look at n_p past states to predict the future n_f data-points by finding other similar instances. We thus measure similarity from taking the hamming distance of each time-series vector. Indeed each vector is in the form $\vec{t}_{vp} = [0, 1, 1, 1, 2, 2, \dots, 2, 1, 0]$, with 0 signifying an away state, 1 a present state and 2 a sleeping state, meaning that taking the Hamming distance between two such vectors is possible. From this a determined number of similar days can be found to use their known future data-points to predict the current future. This is summarized in Algorithm 1.

Data: Time-series vector of past n_p states - t_{vp}

Result: Time-series vector of future n_f expected states - t_{vf}

take current time and give it an index n_n ;

for i index of all past days **do**

set t_{vd} equal to the $n_n - n_p$ to n_n data-points of day_i (i.e the same time range than t_{vp} but in the present);
 store in D_i the Hamming distance of t_{vd} and t_{vp} ;

end

get the indices of the 4 closest days in terms of Hamming Distance from D_i ;

from these selected days get the n_n to $n_n + n_f$ data-points;

get the mode of these selected days and store in t_{vf} ;

return t_{vf}

Algorithm 1: Timetable State prediction


To find the optimal number of days to use to predict future expected states by taking their mode - i.e. should the closest x or y days be used - various prediction scenarios were run with the results presented in section 8.2.2. They show that, for the used dataset (3 months long), the optimal days to take the mode of was 4 as this consistently gave the highest average accuracy and one of the lowest standard deviations⁶.

In deciding how long the vector \vec{t}_{vp} (i.e. the length of time to look back into the past to predict the future

⁶Over 3 months worth of data, about 80 days have their future states predicted and the accuracy of each of these instances is recorded, yielding an average and standard deviation for accuracy over different parameters

states) should be a MATLAB script testing for the accuracy of different settings⁷. In Figure 3.9 the results were plotted and suggest that taking the last 12 hours of house states gives us the most insight into the future - with the accuracy varying little with regards to the time looked into the future, except for very short amounts of time or 12 hours onwards which is because a house state is not likely to imminently change and that as we approach 24 hours we approach a day's length which is typically regular in nature.

Figure 3.9: Model of accuracy using time length past values (sample p) compared to the future hours predicted (sample f)



Data used: University of Massachusetts Amherst SMART dataset [11]*

3.2.4 Machine Learning for Thermostat Setting Prediction

A fundamental part of the functionality of the system is the ability to be able to predict the user's desired thermostat setting given the state of the house and external conditions. This can be achieved through a machine learning algorithm, which regresses through the training data provided by the user to predict a thermostat setting given external conditions not present in the training set. Two regression algorithms were developed, and in this section both are discussed.

Linear Regression

The initial implementation involved implementing a basic linear regression model to fit past data. This works by finding the parameters of the vector $\vec{\theta}$ which forms the prediction function $h(x, \theta) = \sum_{i=0}^n \theta_i x_i = \theta^T x$ where \vec{x} is the input data and n is the number of parameters. To fit the parameters we have a dataset of x 's with it's corresponding values of y . From this a cost function $J(\theta) = \frac{1}{2m} \left(\sum_{i=1}^m (h(x_{(i)}, \theta) - y_{(i)})^2 + \lambda \sum_{i=1}^n \theta_i^2 \right)$ is defined, where m is the number of data-points. Thus by choosing a θ which closely fits the input data we can see that $J(\theta)$ will be minimized. (The $\lambda \sum_{i=1}^n \theta_i^2$ term is there to ensure that θ 's magnitude is not too large - avoiding the problem of the prediction function not being general enough for future, unknown inputs). Minimizing θ is a matter of performing a process called gradient descent which involves following the gradient direction. Figure 3.10 shows 3 arrows each representing the gradient at 3 different points. It can be seen from this that following the gradient leads to a less steep gradient until a minimum is reached - in which case the gradient will be zero⁸.




Figure 3.10: A convex cost function

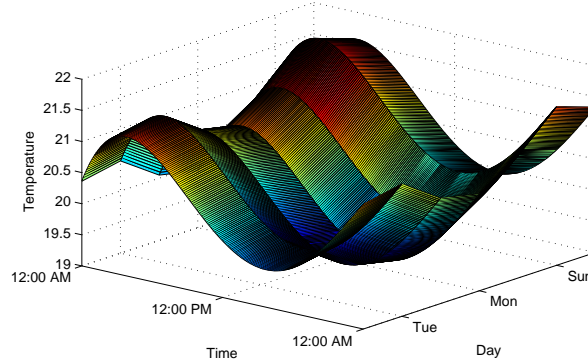
Using the SMART* dataset a θ was found to fit the past data whose features include inside & outside humidity,

⁷This was done by taking the states of each day in a 3 month long dataset separately, removing it from the data set, and trying to predict the future(know to us but unknown to the algorithm) states.

⁸This method also works on non-convex functions but will only return a local minima - however $J(\theta)$ is always convex for the case of linear regression

outside temperature, wind-speed, time of the day and day of the week with output inside temperature. From this Figure 3.11 was made, where all other features except time and day of the week were kept constant, to give an idea of the type of model created. It can be observed that this current model implies a trend of lower temperatures on weekdays with temperature consistently lower during sleeping hours - which implies that the model works in learning from past house activities.

Figure 3.11: Model of a house's temperature with time and day of the week as variables. Gradient descent was used to fit a 4th order function.



Gaussian Process Regression

The issue with using 'non-Bayesian paradigm' algorithms, i.e. ones that typically assume no prior distribution or structure on a hypothesis, is that the regression curve is found (or trained) by minimizing the empirical error. However what is actually desired is to minimize the error on future predictions. This requires a Bayesian setting [15, 22].

Using a Bayesian setting a structure or prior is assumed, which, in the case of GPR, is a Multivariate Gaussian distribution [15, 16]. Parametric techniques, for example linear regression, restrict the class of functions which can fit the training data and then optimize parameters by minimizing the error between the target data set and the predicted values [15, 23]. GPR in contrast allows every possible function but gives each a prior probability, which can then be used to calculate a posterior probability once the training data has been taken into account. In this way then we can think of the GPR as filtering the functions in terms of their likelihood given the training data. As a result, while parametric techniques can generate curves that closely fit the data, only Bayesian techniques like GPR can quantify the expected error and therefore work to minimize it [15, 17, 18, 20].

Notes on Notation:

In this section, bold font indicates vectors and matrices while non-bold indicates scalars and functions. If a character denotes a function then the emboldened version of the same character represents a set or vector of elements taken by evaluating that function at certain points:

- Features vectors are denoted by \mathbf{x} .
- The training data is a set of feature points. In matrix form it is denoted by \mathbf{X} and represents a collection of feature vectors. 'n' is the number of training examples in our training set.
- $f(\mathbf{x})$ (or, in suppressed argument form, f) is the hypothesis function we wish to infer. f^* is the prediction we want to make about the function $f(\mathbf{x})$ evaluated at a point \mathbf{x}^* not present in the training set.
- $y(\mathbf{x})$ (in suppressed argument form as y) is the function of observed thermostat results, i.e. the underlying function we wish to infer plus a noise term: $y(\mathbf{x}) = f(\mathbf{x}) + \epsilon$, where $\epsilon \sim \mathcal{N}(0, \sigma_n^2 \mathbf{I})$ is Independent Additive White Gaussian Noise (IAWGN) with $\text{var}(\epsilon) = \sigma_n^2$. The target data of the training set is denoted by y and is simply a vector of elements of $y(\mathbf{x})$ evaluated at different points.
- ' θ ' Indicates the hyperparameters that are used to define the covariance function, in vector form this is denoted ' $\boldsymbol{\theta}$ '.

It is important to bear in mind that a vector or matrix is a way of expressing a set of data points. If these data points are random then this data set can be written as a random vector or matrix whose elements are random variables each with a probability distribution, which in the case of GPR is Gaussian. As a result data sets, made up of a collection of random variables each with a Gaussian distribution can be expressed as a multivariate Gaussian distribution.

Prior and Hyperparameters:

As previously discussed GPR is a non-parametric technique, which makes it a powerful tool as no assumptions need to be made on the hypothesis function directly, instead only on the probability distribution of the hypothesis

function. The reason that GPR is a good choice in this case is because Gaussian distributions are continuous (important as the hypothesis function we wish to infer is continuous) and flexible. The prior for GPR is defined as:

$$f \sim GP(\mathbf{0}, k(\mathbf{x}_i, \mathbf{x}_j)) \quad (3.3)$$

The above equation states that without any information the distribution of possible hypothesis functions is defined by a Gaussian Process with a mean function of 0 and a covariance function k . As a result to specify the Gaussian prior all that is required is to describe its covariance function (since the mean function can be set to 0 without any loss of generality) [15]. The kernel deployed in this particular application is the squared exponential, also called the Gaussian Kernel, with an added noise term [15]:

$$k(\mathbf{x}_i, \mathbf{x}_j) = \sigma_f^2 \exp\left(-\frac{1}{2}(\mathbf{x}_i - \mathbf{x}_j)^T \mathbf{M} (\mathbf{x}_i - \mathbf{x}_j)\right) + \sigma_n^2 \delta_{ij} \quad (3.4)$$

\mathbf{x}_i and \mathbf{x}_j are input feature vectors, σ_f^2 is the ‘signal power’ which represents the overall scale variation of the latent (not yet manifested) thermostat setting. σ_n^2 is the ‘noise power’ representing the noise variation or ‘jitter’ in the data [23]. Physically this is the manifestation of the users own error in their thermostat setting. For example, on a particularly cold day the user may be tempted to initially turn up their thermostat setting beyond what they actually find comfortable or need. δ_{ij} is the Kronecker product, meaning that the noise term is only added to the variance of a point [23]. This is due to the nature of the noise being IAWGN i.e. the noise on different samples is independent. Finally \mathbf{M} is a diagonal matrix containing the length scales of each feature [15]:

$$\mathbf{M} = \begin{bmatrix} \frac{1}{l_1^2} & 0 & 0 \\ 0 & \ddots & 0 \\ 0 & 0 & \frac{1}{l_n^2} \end{bmatrix} \quad (3.5)$$

Length scales represent the distance in a feature dimension that must be moved before the function value changes significantly. Large length scales imply irrelevant features while short length scales mean that the error in the prediction grows rapidly away from the data points [23]. Having individual length scales for each feature is more flexible than having just one general length scale as there will be variation in length scales from feature to feature.

The error squared exponential kernel was chosen as it provides a smooth (infinitely differentiable) hypothesis and is generally the most widely used and applicable kernel in machine learning.

Hyperparameter Optimization:

To achieve accurate predictions the hyperparameters need to be optimized to ensure the model best fits the training data. In this way the task of ‘teaching’ a Gaussian Process Regressor is therefore equivalent to optimizing the hyperparameters to fit the training data. The marginal likelihood is the ‘model evidence’ and by maximising the marginal likelihood the optimum hyperparameters can be found (See appendix for more information) [15, 18, 21, 22].

The probability of the target data ‘ y ’ given the training data, hyperparameters and assumed prior distribution is the called the likelihood [21, 22]:

$$P(\mathbf{y} | \mathbf{X}, \boldsymbol{\theta}, f) \sim \mathcal{N}(f, \sigma_n^2 \mathbf{I}) \quad (3.6)$$

This distribution makes intuitive sense considering that $y(\mathbf{x}) = f(\mathbf{x}) + \epsilon$ [15, 19] and that the any collection of samples drawn from a Gaussian Process naturally has a Multivariate Gaussian Distribution. The marginal likelihood, $P(\mathbf{y} | \mathbf{X}, \boldsymbol{\theta})$, is found by taking the conditional probability of the likelihood given certain values of f , and then averaging this by integrating over all values of f . This is equivalent to taking the expectation:

$$P(\mathbf{y} | \mathbf{X}, \boldsymbol{\theta}) = E\{P(\mathbf{y} | \mathbf{X}, \boldsymbol{\theta}, f)\} = \int P(\mathbf{y} | \mathbf{X}, \boldsymbol{\theta}, f) P(f | \mathbf{X}, \boldsymbol{\theta}) df \quad (3.7)$$

$P(f | \mathbf{X}, \boldsymbol{\theta}) \sim \mathcal{N}(\mathbf{0}, \mathbf{K})$ is the probability distribution of the function we wish to infer without any target set. In this way, it is equivalent to simply drawing points from our prior distribution. We can see therefore that both the prior set and the likelihood set are both Gaussian. The product of two Gaussians is also Gaussian, integrating over f it can be proved that [15]:

$$\mathbf{y} | \mathbf{X}, \boldsymbol{\theta} \sim \mathcal{N}(\mathbf{0}, \mathbf{K} + \sigma_n^2 \mathbf{I}) \quad (3.8)$$

Referring to the equation for the PDF of a Gaussian Multivariable Distribution, and setting $\mathbf{C} = \mathbf{K} + \sigma_n^2 \mathbf{I}$ the PDF for the Marginal Likelihood (see Appendix) can be defined:

$$P(\mathbf{y} | \mathbf{X}, \boldsymbol{\theta}) = 2\pi^{-(\frac{n}{2})} |\mathbf{C}|^{-0.5} \exp\left(-\frac{1}{2} \mathbf{y}^T \mathbf{C}^{-1} \mathbf{y}\right) \quad (3.9)$$

In its current form it is not particularly conducive to the task of optimization. A common and successful

technique for optimizing any parameters is to turn the problem into a minimization task, and find a global or acceptably good local minima using gradient descent. Taking the natural log of the equation:

$$\log(P(\mathbf{y}|\mathbf{X}, \boldsymbol{\theta})) = -\left(\frac{n}{2}\right) \log(2\pi) - \frac{1}{2} \log(|C|) - \frac{1}{2} \mathbf{y}^T C^{-1} \mathbf{y} \quad (3.10)$$

Rearranging and multiplying by negative one we end up with what is known as the 'Negative Log Marginal Likelihood' or NLML [21, 22]:

$$-\log(P(\mathbf{y}|\mathbf{X}, \boldsymbol{\theta})) = \frac{1}{2} \mathbf{y}^T C^{-1} \mathbf{y} + \frac{1}{2} \log(|C|) + \frac{n}{2} \log(2\pi) \quad (3.11)$$

Note that by minimizing the NLML we are maximising the Marginal Likelihood and therefore gradient descent techniques can be applied. First the hyperparameters are initialised randomly and then taking the partial differential of the NLML with respect to each hyperparameter the error curve can be traversed (error curve is the NLML as a function of the hyperparameters) to find a minima. Denoting the NLML as $L(\boldsymbol{\theta})$ [22]:

$$\frac{\partial L(\boldsymbol{\theta})}{\partial \theta_i} = \frac{1}{2} \text{tr} \left(C^{-1} \frac{\partial C}{\partial \theta_i} \right) - \frac{1}{2} \mathbf{y}^T C^{-1} \frac{\partial C}{\partial \theta_i} C^{-1} \mathbf{y} \quad (3.12)$$

A basic gradient descent function can be used to update the hyperparameters by subtracting a value proportional to the partial derivative of the NLML with the hyperparameter in question. Over the process of many iterations, and given an appropriately sized step size α (i.e. one that is not too large so as to diverge or too small so as to converge too slowly), the NLML will converge to a point of zero gradient. As we are subtracting this term this convergence limits towards a minima, not a saddle point or maxima:

$$\theta_i = \theta_i - \alpha \frac{\partial L(\boldsymbol{\theta})}{\partial \theta_i} \quad (3.13)$$

Unfortunately this problem is not convex (i.e. the error curve is not convex) and therefore there does not exist a global minima. Instead there is the potential for many minima and maxima depending on our selection of training data. As a result all that can be hoped for is that gradient descent gives us an acceptable local minimum.

The convergence point depends on what the initialisation of the hyperparameters; different initialisations lead to different starting points on the error curve and therefore potentially different local minima. The practical approach taken to solve this issue was to perform a grid search over an area of the error curve (that seemed most logical given our instincts on the data), perform gradient at points initialised in divisions of this area and then choose the area that seemed to give the lowest average NLML. The division with the lowest NLML can then be selected as the new area of consideration, and the process can be repeated till an appreciable initialisation range is found. This, and the computational complexity of all the algorithms discussed in this section, is discussed in more detail in the results section.

Intuition behind Gaussian Process Regression:

With a Gaussian prior the assumption is made that the space of possible functions that we can use as our hypothesis is normally distributed. By then observing data, i.e. getting a training set, the distribution can be conditioned on the training set. By then finding the mean or average of this distribution function that best, or most likely fits the user's desired thermostat setting given certain environmental conditions, can be found. Points in the training data set can be seen as anchors in this function distribution having zero variance [15, 17, 20, 23].

Considering a new input feature vector \mathbf{x}^* and a desired prediction about $f(\mathbf{x}^*)$ (f^* for convenience) then adding this point to the training data results in a joint Gaussian distribution [15]:

$$\begin{bmatrix} \mathbf{y} \\ \mathbf{f}^* \end{bmatrix} \sim \mathcal{N} \left(\mathbf{0}, \begin{bmatrix} \mathbf{K} + \sigma_n^2 \mathbf{I} & \mathbf{K}_* \\ \mathbf{K}_*^T & \mathbf{K}_{**} \end{bmatrix} \right) \quad (3.14)$$

Here $\mathbf{K}_* = \begin{bmatrix} k(\mathbf{x}_*, \mathbf{x}_1) \\ \vdots \\ k(\mathbf{x}_*, \mathbf{x}_n) \end{bmatrix}$ denotes the covariance matrix (or in this case, as \mathbf{x}^* is a single vector, a vector) of the input feature vector with each of the training examples. As already discussed the objective is to determine the best or most likely value of f^* given the training data. As a result what is of greatest relevance is the conditional probability of the predicted value given the training data. This is known as the 'posterior probability', and can be derived from Bayes Theorem:

$$P(f^*|\mathbf{X}, \mathbf{y}, \boldsymbol{\theta}) = \frac{P(\mathbf{f}|\mathbf{X}, \boldsymbol{\theta}) P(\mathbf{y}|\mathbf{X}, \boldsymbol{\theta}, \mathbf{f})}{P(\mathbf{y}|\mathbf{X}, \boldsymbol{\theta})} \quad (3.15)$$

Referring to the set of definitions given above we can write this alternatively as [15]:

$$\text{posterior} = \frac{\text{prior} \times \text{likelihood}}{\text{marginal likelihood}} \quad (3.16)$$

The marginal likelihood is a 'normalizing constant', i.e. a constant as it is the Likelihood function marginalized over f . Therefore the posterior probability can be more simply written as [15, 21, 22]:

$$P(f^* | \mathbf{X}, \mathbf{y}, \boldsymbol{\theta}) \propto P(\mathbf{f} | \mathbf{X}, \boldsymbol{\theta}) P(\mathbf{y} | \mathbf{X}, \boldsymbol{\theta}, \mathbf{f}) \quad (3.17)$$

As the prior and the likelihood function both follow a Gaussian distribution then the posterior probability distribution is also Gaussian (intuitively multiplying two exponentials together renders another exponential). Using some basic results found in Linear Algebra [19] this Gaussian distribution can be shown to have a mean and standard deviation of the following [15]:

$$P(f^* | \mathbf{y}, \mathbf{X}) \sim \mathcal{N}(\mathbf{K}_*^T (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{y}, \mathbf{K}_{**} - \mathbf{K}_*^T (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{K}_*) \quad (3.18)$$

The above equation describes the distribution of the possible hypothesis functions given the data available. The average or mean of this distribution gives the most likely regression function given the data; therefore the mean of this distribution is the best estimate of f^* :

$$\hat{f}^* = \mathbf{K}_*^T (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{y} \quad (3.19)$$

The confidence in the prediction is described by the variance $\mathbf{K}_{**} - \mathbf{K}_*^T (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{K}_*$, as a result GPR can calculate a best estimate of the desired thermostat setting given the training data and also gives a metric for the confidence the system has in the prediction.

3.3 User Interface

3.3.1 Website

The website was designed to be easy to use and minimalistic as to keep it as user friendly as possible. All text is placed inside lightened boxes, otherwise it became difficult to read against the background.

Trying to access any page other than the login page before logging in, will result in a redirect to the login page. Upon logging in, the user is forwarded to the main home page. The home page greets the user by their username, giving the site a more personal feel. This page also contains a toggle slider to turn the system off, for instance if they are going away for an extended period of time.

The control page gives direct control of the lights remotely via a simple toggle switch. The switch also updates in real time, so if the light is turned on or off manually, this will instantly be updated on the website. This means that you can check if you have left your light on when you are not at home and even turn it off if you have.

The statistics page primarily displays the the schedule for the house, showing the user when they are expected to be in, out or asleep. If the user disagrees with the prediction, they are also able to update the schedule themselves, with user corrections being displayed in a different colour.

At the bottom of the page, the user can view graphs of past external weather conditions, as well as past internal conditions, including the power usage of the house.

The thermostat page has the primary use of updating the current thermostat setting manually. Every time the user does this, the value is put into a table along with the most recent external and internal conditions, so that the algorithm that selects the temperature can learn from this.

The user can also change the maximum and minimum temperatures that the algorithm selects, as well as their preferred sleeping temperature.

3.3.2 Mobile App

Along with the website the decision was made to develop an application (or app) for smartphones, this was to be done for the android operating system and then would later be ported to the it's various competitors. The reasons for this were, that android was the most popular operating system for smartphones; the multitude of available examples and developer information, and most importantly the open (hence free) source nature of the operating system.

Using Eclipse combined with android SDK (software development kit) tools a basic app was built with a separate temperature and lights page which would then be used to change the thermostat setting and turn the lights on and off (with toggle buttons). Using an open source android websocket library the app is able to communicate with the server and adjust the various attributes of the system.

Section 4

System Analysis and Use

4.1 Savings Performance

Justifying the home automation system in terms of saving energy over the conventional thermostat is critical in justifying its usefulness. Firstly, the motivation to cut down the use of domestic heating is significant; housing accounted for over 28.3% of the UK's energy consumption in 2009[1] at 501 TWh. Of this, a staggering 65.7% is accounted for the heating alone. Naturally, cutting this will make a significant impact on total energy usage and the environment.

In order to estimate the energy the system can save, it is important to first investigate the usage from conventional thermostats. In recent years, a continual improvement in minimising heat loss is shown the graph below. In 2008, the total heat loss is at 253.7 W/K. This means that for every degree of temperature difference between inside and outside, over 250 Watts are needed to maintain the desired internal temperature.

Figure 4.1: Energy loss by house insulation





Figure 4.2: Year-on-year winter internal to external temperatures



From this data, an estimation of the power usage can be made from the average winter temperature difference between inside and out. Taking data from 2008, the average temperature difference was 10.9 °C. With a thermal loss of 253.7 W/K, this means a power of 2.77 kW is required to heat the average home. The weighted average fuel price is 6.02p/kWh: meaning it costs 16.7 pence for every hour the central heating is switched on. Assuming that the heating is used for 12 hours a day, this means a daily cost of £2. Considering that the weekly expenditure on lights, heating and power is £18.90, this seems to be a reasonable estimate.

Above are the approximate hourly savings the home automation system can provide the average UK household during winter months per hour. Although these appear small, if the heating was turned down for 7 hours a day by 2 °C whilst the occupants are sleeping, and 1 hour or so of unnecessary heating is reduced every day, then the daily

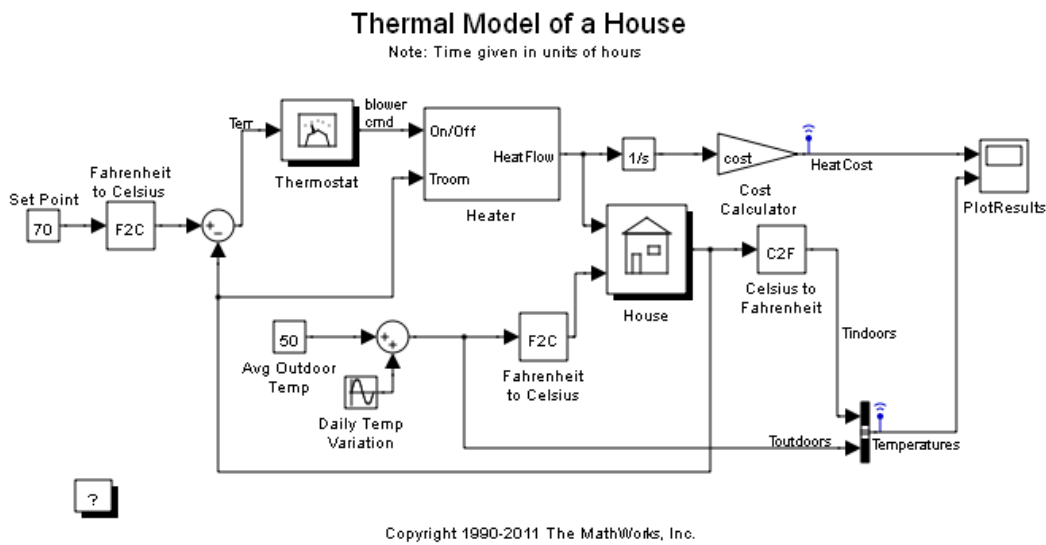
System prediction benefit	Approximate hourly saving (£)
Turning down temperature in sleep state	1.53 p/K
Turn off heating whilst house is unoccupied	16.7 p

saving the system would provide will be around £0.40. This is a reduction in 20%, which over the course of a year could be £130 for the average UK gas bill of £653 on direct debit.

Another way of estimating the amount the system could save is by thermally modelling a house and calculating the energy used for different thermostat settings: one set like a traditional thermostat, which maintains a temperature for 12 hours a day through the week and the entirety of Sunday; and one using simulated occupation data, which a machine learning algorithm could hopefully predict.


Fortunately, MATLAB has a Simulink model for a house heating system, so making the simulations was reasonably straightforward. The program was modified by firstly matching the program parameters to the thermal resistance, temperature and average p/kWh of a UK home as specified above. Additionally, the constant thermostat setting was replaced by a switchable signal generator giving either the AutoHome or conventional thermostat settings. The schematic for this model is shown in Figure 4.3.

Figure 4.3: The House Heating Simulink Model




The model works by switching on/off a "blower", which heats 1 kg of air every second to 50 °C when the internal temperature of the house deviates from the setpoint by more than 0.5 °C. A feedback loop inside the house block takes the thermal losses from the temperature gradient into account. Integrating the thermal power going into the house gives the energy used, which the cost can be calculated from. A weeks worth of thermostat settings (168 hours) is fed into the system for the conventional and occupancy setpoints, and the results shown in the graphs below. The external temperature, shown as green in the lower plots, oscillates sinusoidally with an amplitude of 4 °C around the average winter temperature in 2008 of 6.4 °C.

Figure 4.4: Simulation result for a week's worth of conventional thermostat program heating cost and temperatures



The first thing to note is that the estimated costs are in the same ballpark of those predicted by simple calculation earlier at £21.90 and £19.48 respectively, which represents a saving of almost £2.50. Although this is a quite a modest saving, it is worth revisiting the other aim of the automation system, which is to increase the comfort and convenience of the house's heating. Clearly the occupancy of the house is non uniform unlike the setting of the conventional system, which means there will be times when the house is either unnecessarily heated or at an uncomfortable temperature.

Figure 4.5: Simulation result for a week’s worth of predicted occupancy and corresponding thermostat setting cost and temperatures



4.2 Hardware Performance

Each module of the hardware was first developed separately and tested so in order to fix any issues at their earliest stage. These were then all connected and the programmes combined to form a singular master programme which would control the whole mbed system.

First the digital temperature sensor was set up and programmed as discussed in section 3.1.3 after which initial tests found the temperature to be very low, however with the addition of pull-up resistors on the SCL and SDA pins the temperature was found to be as expected, this was then tested using an advanced handheld temperature sensor and was found to be accurate to $\pm 0.5^{\circ}\text{C}$ as expected.

With the temperature sensor now completed the humidity sensor could be utilised, as it requires the local temperature in order to calculate the correct humidity from the input voltage which is measured using an ADC port on the mbed. This was then tested by measuring the outside environment with the humidity returned from the weather API, this was found to return similar values though the humidity sensor itself was not particularly stable and varied by several % points. However this would be as expected as the datasheet tells us the $\pm 3\%$.

With the local system complete, the wifly module was set up and connected to an access point setup using a laptop (as it cannot connect to enterprise enabled wifi). It was also tested using a standard home router which was supplied by the ISP company. After which a websocket connection was set up, this was generally successful but ran into a number of problems.

The first and major problem was that the websocket would close after an unknown period of time for unknown reasons. Furthermore the `is_connected` function failed to report when the websocket would disconnect, the general solution came from the common ping-pong idea. Where a ping was regularly sent from the mbed and the server would then return pong, if this return 'pong' was not detected by the mbed mode it would then re-join the websocket, however this was only partially successful and sometimes failed completely, and so a reset then reboot command was sent to the wifly module after which it would then reconnect quickly with no failures.

4.3 Software Performance

4.3.1 Timetable Prediction

Every three hours¹ the timetable prediction program (described in section 3.2.3) is called to predict what it expects the current day to look like, as well as what it expects (though with lower confidence) the other days of the week. This is shown in Figure 4.6 which shows that the user is allowed to override the estimated timetable for advanced energy savings - however as current systems work purely by such scheduling and are often seen as inefficient this use is discouraged.

This timetable is used if the current state is either **away** or **asleep** and is expected to be **inside** (which also implies awake) in the next 15 minutes². When this is the case the house is heated up to the expected comfortable temperature (determined by the algorithm in section 3.2.4). It will sometimes be the case that the user does not come back when expected. In such a case a half-hour window will be given for the house state to become **inside** (i.e. detect someone is present) before ignoring this timetable and instead waiting for the house state to actually become **inside** to avoid energy waste (at the expense of compromising user comfort).

Testing the real world use of the predictor was a challenge as a lot of data would need to be collected - a feat not possible in the amount of time given. Instead the UMass SMART* dataset[11] was used to test this data and determine the accuracy (which also allowed for finding the best parameters) as seen in Figure 3.9. From this we can see that the expected accuracy is 85% which is a reasonably good performance which however requires a fallback

¹This value can be changed but was chosen to be three hours as the algorithm works best by regressing on 6 hours of past data - however to accommodate for unexpected events and offer more granularity 3 hours was chosen as a good compromise

²The certainty - i.e. if the algorithm is 50% sure or 75% sure - can be changed but is set to 50 by default


& override in case of failure - which is described above.

Figure 4.6: Timetable presented on website

Legend:

Red: Asleep - Green: Inside - Blue: Away

Dark: Algorithm Predicted State - Light: User Set State



4.3.2 Website Usage

The website was primarily tested in college using Google Chrome. This is because the site is only accessible internally to college (although it can be connect to elsewhere using the VPN) and Chrome is the only browser on the college computers that is up to date. Using a recent browser is important to the website, as it has regular use of websockets; a protocol only introduced in HTML5, along with a few other features not supported in older browsers. Unfortunately the other browsers on campus are not up to date enough to support websockets.

The first test for each page was to simply check for errors in the console window and address any issues that arose. After this, any major change made to the website was heavily tested for bugs by a handful of people. Any changes that were needed were made, with this process being repeated again for the updated version. The websockets could be tested easily by checking the websocket server log for messages, with the results of anything returned being displayed directly on the website.

Ajax requests could be checked in one of two ways. Firstly it can be checked that they are actually called and receive a response by using the network console window. This lets you see detailed information about any request, such as what data is sent with the request and any data that is returned. Secondly, because the ajax requests all deal with database handling, it is possible to check the relevant tables using mysql directly.

It is possible to check that the weather API script is collecting data correctly by viewing the graphs on the statistics page.

Once the website was fully functional in Chrome, it was also tested in latest versions of Internet Explorer and Firefox. A few small modifications were needed to ensure cross browser compatibility which were easily made. The site was also tested on IE9 mobile, where it failed due to lack of support for websockets, but was otherwise fine and also on safari for iOS, where it worked fully. The site is not however optimised for mobile use, this is where a phone app would come in use.

4.4 System Practicality

The practicality of the system was determined by assessing the ease at which it could be installed into a house. It is important the the entire system can be easily retro fitted to a house. Otherwise, only the relatively small number of new builds could get the benefits the system has to offer, which would render it both economically invalid and removes any real potential to impact the national or worldwide energy consumption. For each of the hardware modules, important considerations are how the component will be installed, maintained, the intrusion caused to the users and the security of it.

The majority of UK households have a gas fired boiler, which is controlled by either bimetallic coil switch in mechanical thermostats or a relay in programmable flavours. For these, only four wires are needed: live, neutral and the two switch wires for the relay. Connecting the main mbed module to this can either be done directly or with a separate boiler relay unit included connected to the wireless ZigBee network. The thermostat then maintains the temperature at a given setpoint by switching on/off this relay. Whether the house has an oil, electric or gas boiler or a ground source heat pump, the thermostat can control a local temperature with ignorance. Installing such a system will be identical to any commercial programmable thermostat, which can be done by any plumber or even by the homeowner with relative ease.

Similarly, with the light switches, the electronics is simple enough to be integrated into a conventional light switch housing, and can be replaced easily. The power supply can be derived off the mains that it is switching.

Furthermore, since the lightswitch is effectively just a wirelessly controllable mains supply with a bulb attached to it, the design can be used in any domestic appliance and embedded into a traditional mains socket or integrated into the appliance itself.

As for the power meter, the same current clip design is used in many of the commercially available ‘smart meters’ that are designed for the user to simply install themselves due to the safety of not having to touch the mains at any point. Although not implemented on the prototype, a very simple and sensible way to power to unit is through the voltage sense transformer by rectifying, smoothing and applying to the regulation circuit.

Motion sensors need to be in critical places such as main living spaces such as kitchen, bathroom and lounges, but also in bedrooms to determine sleep state. Naturally, the idea of having a wirelessly transmitting motion sensor especially in a bedroom may seem like an invasion of privacy, in reality they are only a passive infrared receiver with a digital output giving a change in radiation due to motion. Therefore, no information apart from the activity of the house, which is what the system is trying to ascertain, can be extrapolated.

4.5 System Unit Cost

Note: A breakdown of costs can be found in the appendix section 8.3.

As can be seen in the table above the final cost of the prototype came to £212.34 (including breadboard) . While this is above our expected costs this does not take into account any benefits from economies of scale, plus the advantages from component switching when working on PCB (as opposed to breakout boards required for breadboard). Furthermore the mbed development board is the most expensive unit for our prototype, however when finally developed only the microprocessor (and some other debugging components) will be required. This would drop the cost of the microprocessor from £41.38 to just £5.51 (for a single unit, lower for multiples) and so already it can be seen that there is significant scope for reductions before the final design.

Section 5

Further Development

There are many ways in which this system could be enhanced, for example adding cameras or microphones to increase the detection rate of human presence. However many of these come at the cost of overcomplicating the device (and its installation) - for example the proposition of adding a camera not only complicates the installation and adds little benefits but also introduces privacy problems (many people would object to the installation of a camera in their home) - ultimately causing the product to be less attractive. Thus further developments must take this into account to genuinely improve the product and make it more attractive and useful rather than loading it with features.

5.1 LCD Screen

5.2 Modular Network Setup

One major improvement to our design would be the addition of a simple, ease of use module setup feature. This would be where different components could be purchased separately by the consumer and then with a simple few button touches would be connected into the network as a fully functional device, this could be done by using some of the abilities of Xbees which allows new Xbees to be easily included into a new network using an API framework.

5.3 Smart Door Lock

This improvement would involve offering a module which can be installed on the main door of the house and compliment a usual key lock by offering a wireless solution. The advantages of such a lock would be the ability to open the main door simply via a digital key adding more possibilities for key control. For example the keys could be given to guests temporarily or only work for the cleaning lady at certain times and deactivated remotely. Digital security would evidently be a high priority of such a device.

Additionally this type of device, while very new, is not innovative in itself as there are already two companies trying to enter this market - [Lockitron](#) and [August Lock](#). The real interest in offering this as a module would be near 95-99% certainty in house occupancy¹. Indeed by combining the house motion sensors with main door traffic flow (which can be determined whether someone has just entered the house or not). The state can be determined by a process similar to:

- **Key used to enter house:** Assume someone is entering house thus presence is inferred
- **Key not used:** Assume that someone is either leaving the house or guests are being let in. Differentiate by:
No motion detected in house up to 5 minutes after event: Assume house empty (otherwise motion would be detected when moving inside house)
Motion detected 5 minutes after event: Infer presence in house

While a Hidden Markov Model will still need to be relied upon to clearly differentiate states this method is expected to greatly increase house estimation accuracy while adding a useful module to the system.

¹Assuming no pets which is another issue to deal with

Section 6

Conclusion

Section 7

References

- [1] Jason Palmer, Ian Cooper *Great Britain's Housing Energy Fact File* https://www.gov.uk/government/uploads/system/uploads/attachment_data/file/48195/3224-great-britains-housing-energy-fact-file-2011.pdf
- [2] L.D.Shorrock, J.I.Utley *Domestic Energy Fact File 2008* http://www.bre.co.uk/filelibrary/pdf/rpts/fact_file_2008.pdf 2008.
- [3] DOE *Annual Report on Fuel Poverty Statistics 2013* https://www.gov.uk/government/uploads/system/uploads/attachment_data/file/199833/Fuel_Poverty_Report_2013_FINALv2.pdf 2013.
- [4] MarketsandMarkets *Global Home Automation & Control Systems Market is expected to reach \$35,627.83 million by 2016* <http://www.marketsandmarkets.com/PressReleases/home-automation-control-systems.asp> September 2011.
- [5] MarketsandMarkets *Internet of Things (IoT) & Machine-To-Machine (M2M) Communication Market* <http://www.marketsandmarkets.com/Market-Reports/internet-of-things-market-573.html> September 2012.
- [6] Farnell UK *NXP - LPC1768FBD100 - MCU, 32BIT, ARM CORTEX M3, 100LQFP* <http://uk.farnell.com/nxp/lpc1768fbd100/mcu-32bit-arm-cortex-m3-100lqfp/dp/1718549MPKG> June 2013.
- [7] Mike Elgan *Home Tech: There's an Easier, Affordable Future for Home Automation* <http://www.houzz.com/ideabooks/4307872/list/Home-Tech--There-s-an-Easier--Affordable-Future-for-Home-Automation> September 2012.
- [8] HM Government *Helping households to cut their energy bills* <https://www.gov.uk/government/policies/helping-households-to-cut-their-energy-bills> May 2013.
- [9] Department of Energy and Climate Change *Smart Metering Implementation Programme Consultation document* https://www.gov.uk/government/uploads/system/uploads/attachment_data/file/42953/6129-consultation-second-version-smets.pdf August 2012.
- [10] Hart, G. W. *Residential energy monitoring and computerized surveillance via utility power flows* <http://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=31557> 1989.
- [11] Sean Barker, Aditya Mishra, David Irwin, Emmanuel Cecchet, and Prashant Shenoy *Smart*: An Open Data Set and Tools for Enabling Research in Sustainable Homes* <http://lass.cs.umass.edu/papers/pdf/sustkdd-smart.pdf>
- [12] Department of Energy *Energy Efficiency and Renewable Energy* http://energy.gov/sites/prod/files/energy_savers.pdf 2011.
- [13] Washington State University *WSU CASAS Datasets* <http://ailab.wsu.edu/casas/datasets/index.html>
- [14] James Scott, A.J. Bernheim Brush, John Krumm, Brian Meyers, Mike Hazas, Steve Hodges, Nicolas Villar *Pre-Heat: Controlling Home Heating Using Occupancy Prediction* <http://research.microsoft.com/en-us/um/people/jckrumm/publications%202011/preheat-ubicomp2011%20-%20distribute.pdf> September 2011.
- [15] Carl Edward Rasmussen and Christopher K. I. Williams *Gaussian Processes for Machine Learning* <http://www.gaussianprocess.org/gpml/chapters/> 2006.
- [16] Jonathan Taylor *Statistics 191: Introduction to Applied Statistics* <http://www-stat.stanford.edu/~jtaylo/courses/stats191/notes/simple.pdf> January 2009.

- [17] Carl Edward Rasmussen, Hannes Nickisch *Documentation for GPML* <http://www.gaussianprocess.org/gpml/code/matlab/doc/> January 2013.
- [18] Manuel Blum, Martin Riedmiller *Optimization of Gaussian Process Hyperparameters using Rprop* http://ml.informatik.uni-freiburg.de/_media/publications/blumesann2013.pdf 2013.
- [19] University of Manchester *Multivariate Normal Distribution* [http://www.maths.manchester.ac.uk/~mkt/MT3732%20\(MVA\)/Notes/MVA_Section3.pdf](http://www.maths.manchester.ac.uk/~mkt/MT3732%20(MVA)/Notes/MVA_Section3.pdf) 2010.
- [20] Pedregosa et al. *Scikit-learn: Machine Learning in Python: Gaussian Processes* http://amueller.github.io/modules/gaussian_process.html 2013.
- [21] Christopher K. I. Williams, Carl Edward Rasmussen *Gaussian Processes for Regression* <http://mlg.eng.cam.ac.uk/pub/pdf/WilRas96.pdf> 1996.
- [22] Ed Snelson *Tutorial: Gaussian process models* <http://mlg.eng.cam.ac.uk/tutorials/06/es.pdf> October 2006.
- [23] David P. Williams *Assessing Approximations for Gaussian Process Classification* <http://people.ee.duke.edu/~lcarin/David1.27.06.pdf> January 2006.
- [24] ARM NXP *32-bit ARM Cortex-M3 microcontroller Rev. 9 — 10 August 2012 Product data sheet* http://www.nxp.com/documents/data_sheet/LPC1769_68_67_66_65_64_63.pdf August 2012.
- [25] Strategy Analytics *A Quarter of Households Worldwide Now Have Wireless Home Networks* <http://www.strategyanalytics.com/default.aspx?mod=pressreleaseviewer&a0=5193> April 2012.

Section 8

Appendix

8.1 Hardware

8.1.1 mbed

Main Control Code

```
1 #include "mbed.h"
2 #include "TMP175.h"
3 #include "WiflyInterface.h"
4 #include "HTTPClient.h"
5 #include "Websocket.h"
6 #include "SDFileSystem.h"
7 #include "picojson.h"
8 #include "header.h"
9 #include "xbee.h"
10 #include "MODSERIAL.h"
11 #define ATIClength 23 // 23 gives best results
12
13
14 AnalogIn ain(p18); // Humidity pin
15 Serial wifi(p13,p14); // Wifly module pins
16 DigitalOut rst1(p26); // Temperature reset pin
17 DigitalOut WiFlyreset(p30); //direct reset signal into WiFly for errorhandler
18 WiflyInterface wifly(p13, p14,p19,p20,"TP-LINK_804327", "C1804327", WPA); //wifly module set-up
19 SDFileSystem sd(p5, p6, p7, p8, "sd"); // SD card set up
20
21 //Tickers set for recurring functions
22 Ticker sendPower;
23 Ticker measureMotion;
24 Ticker lightstatus;
25 Ticker dataSend;
26
27 TMP175 mySensor(p28,p27); // Temperature set up
28 Websocket wsoc("ws://ee-ug1.ee.ic.ac.uk:8888/ws"); //websocket address
29 picojson::value v; //JSON parsing variable
30 XBEE myxbee; //Xbee function
31 MODSERIAL pc(USBTX, USBRX); //Debugging using serial terminal
32 MODSERIAL xbee(p9,p10); //Modserial allows increased buffer size for xbee
33 DigitalOut led(LED1);
34 DigitalOut flash(LED4);
35
36 //initialise variables
37 char response[ATIClength];
38 char orderedresponse[ATIClength];
39 char str[512];
40 char recv[512];
41 float humidity;
42 float Temp;
43 float thermostat = 15;
44 int pCurrent=0;
45 int pPhase=0;
46 int pVoltage=0;
47 int motionCount =0;
48 int bufferpos= 0;
49 int a=0;
50 //64 bit addresses of XBEEs
51 char MR0address[8]= {0x00,0x13,0xA2,0x00,0x40,0x8B,0x71,0x55}; //motion sensor XBEE
52 char MR2address[8]= {0x00,0x13,0xA2,0x00,0x40,0x8B,0x70,0x79}; //power monitoring XBEE
53 char MR3address[8]= {0x00,0x13,0xA2,0x00,0x40,0x8B,0x71,0x45}; //Light control XBEE
54 //extern char MR1address[8];
55 extern char * test;
56 extern char * MR1address; //pointer to array which holds the currently selected address.
57 bool lightStatus=false;
58
59
60 void swapaddress(char address[])
61 {
62 /* used to swap xbee address as required for communication with the
```

```

63     different xbees */
64
65     MRIaddress=&address[0];
66
67 }
68
69 void readATICresponseISR(MODSERIAL_IRQ_INFO *q)
70 {
71     MODSERIAL * serial= q->serial;
72
73
74
75     if(serial->readable())
76         response[bufferpos]=serial->rxGetLastChar();//read in char but dont delete from circular buffer
77
78     bufferpos++;
79     pc.printf("IRQ REPLAY %d IS %2X LONG\n",bufferpos, response[bufferpos]);
80     if(bufferpos == ATIClength) {
81         bufferpos = 0;
82     }
83     if(serial->rxBufferFull()==true) {
84         serial->rxBufferFlush();//flush the buffer once it is full.
85         if(serial->rxBufferEmpty()==true)
86             pc.printf("BUFFER FLUSHED\n");
87         bufferpos=0;
88     }
89 }
90
91
92 void readATICresponseISR2()
93 {
94
95     int pos;
96     if(xbee.readable()) {
97         response[bufferpos]=xbee.getc();//clear bytes from FIFO
98     }
99
100     bufferpos++;
101     pc.printf(" %d IS %2X LONG\n",bufferpos, response[bufferpos]);
102
103     if(bufferpos == ATIClength-1) {
104         bufferpos = 0;
105         pos=myxbee.lastpacketpos(response,ATIClength);
106
107         pc.printf("packet starts here at byte %d pos\n the packet delimiter is %2X \n",pos+1,response[pos]);
108         pc.printf("DIGITAL HEX VARS ARE %2X and %2X\n",response[ATIClength-4],response[ATIClength-3]);//extract the
109         digital values from sent packet
110     }
111 }
112
113 }
114
115 void lightCheck()
116 {
117     //function to call xbee light check function and to send the status to the sever
118
119     __disable_irq();
120     myxbee.SENDLIGHTSTATUSREQ();
121     //sprintf(str, "{\"type\":\"light_request_response\" , \"light_ID\":1 , \"status\": %d }", lightStatus);
122     //wsoc.send(str);
123     __enable_irq();
124 }
125
126 void pingpong()
127 {
128     /* ping pong function used to detect whether websockets are still running
129     and then reset and reboot wifly module then reconnect to websocket */
130
131
132     __disable_irq();
133
134     pc.printf("is connected %d\n",wifly.is_connected());
135     wsoc.send("{\"type\":\"ping\"}"); //send ping tp sever
136     wait_ms(100);
137     //if server responds then websockets are working so return//
138     if(wsoc.read(recv)) {
139         pc.printf("%s\n",recv);
140         //if(strcmp("{\"type\":\"pong\"}",recv) == 0)
141         __enable_irq();
142         return;
143     }
144     //otherwise reset and reboot wifly module
145     WiFlyreset=0;
146     pc.printf("RESETTING WIFLY\n");
147     wait_ms(200);//make an active low pulse
148     WiFlyreset=1;
149     wait_ms(100);//give time to reset
150     wifly.send("reboot",6); //send reboot cmd to wifly module
151     wait_ms(200);
152     wifly.init(); //Use DHCP
153     wait(0.5);
154     wifly.connect(); //connect to wifi

```

```

155     wait(0.5);
156
157     //websocket reconnect
158
159     if(wsoc.connect() == 1) {
160         pc.printf("This was a triumph!");
161     } else {
162         pc.printf("failed!");
163         // pingpong();
164     }
165     __enable_irq();
166     return;
167 }
168
169
170
171 int main()
172 {
173     //set baud rate
174     pc.baud(115200);
175     xbee.baud(115200);
176
177     //wifly reset at start
178     __disable_irq();
179     WiFlyreset=1;
180     wait(2);
181     WiFlyreset=0;
182     wait(2);
183     WiFlyreset=1;
184
185     //set xbee addresses
186     MR1address=&MR2address[0];
187
188
189     pc.printf("Test Wifly!\r\n");
190
191
192     mySensor.vSetConfigurationTMP175(SHUTDOWN_MODE_OFF|COMPARATOR_MODE|POLARITY_0|FAULT_QUEUE_6|RESOLUTION_9,0x48); //
193     Temperature set-up
194     wait_ms(400);
195     wifly.reboot(); //reboot wifly
196     wait(1.0);
197     wifly.init(); //Use DHCP
198     wait(1.0);
199     wifly.connect(); //connect to wifi
200     wsoc.connect(); //connect to websocket
201     __enable_irq();
202
203
204     sendPower.attach(&myxbee,&XBEE::SENDPOWERREQ, 5);
205     //ticker polling power sensor for data
206     wait_ms(1000);
207
208     measureMotion.attach(&myxbee,&XBEE::SENDMOTIONREQ,3.0);
209     //ticker polling motion sensor
210     wait_ms(1000);
211
212     lightstatus.attach(&lightCheck,3);
213     //ticker polling whether the light is on or off
214
215     dataSend.attach(&senddata,17);
216     //send power, humidity, temperature and motion data to server
217
218     while (1) {
219
220
221         //*****Temp sensor*****//
222
223         Temp=mySensor; //get temperature in variable 'temp'
224
225         //*****Humidity Sensor*****//
226
227         humidity = ain.read();
228         humidity = (((humidity*5)/3.3)-0.1515)/0.00636;
229         humidity = humidity/(1.0546-(0.00216*Temp));
230
231
232
233         //*****recieve section*****//
234
235         recv[0] = '\0'; //set recieve string to 0
236         wait_ms(1500);
237
238         __disable_irq();
239         if(wsoc.read(recv)) { //if communication recieve pass data to recieve function
240
241             pc.printf("%s\r\n", recv);
242             recievedata();
243
244         }
245         else {
246             //otherwise check websocket is still running

```

```

247         pingpong();
248     }
249     __enable_irq();
250
251     //end of continious while loop
252 }
253
254 }
255

```

code/homeauto-main.cpp

```

1  #ifndef header_h
2  #define header_h
3  #include "Websocket.h"
4
5  #include "picojson.h"
6  #include "mbed.h"
7  #include "xbee.h"
8
9
10
11
12
13
14  extern Websocket wsoc; //websocket address
15  extern picojson::value v;
16
17  extern char str[512];
18  extern char recv[512];
19  extern float humidity;
20  extern float Temp;
21  extern float power;
22
23  extern float thermostat;
24  extern XBEE myxbee;
25
26  void lightsignal();
27  void senddata();
28  void thermoupdate();
29  void recievedata();
30  void powerconv();
31
32  #endif

```

code/homeauto-header.h

```

1  #include "header.h"
2  #define PI 3.1415926
3
4  extern void swapaddress(char address[]);
5
6  extern char MR0address[8];
7  extern char MR2address[8];
8  extern bool lightStatus;
9
10 float PCurrent;
11 float PVoltage;
12 float PPhase;
13 float power;
14
15
16 void senddata()
17 {
18     /* function sends temperature, humidity, power and motion data back to the server at regular intervals */
19     __disable_irq();
20     powerconv(); //function to work out correct power measurement from calibration data
21
22     sprintf(str, "{\"type\":\"house_measurement\" , \"Temperature\":%f, \"Humidity\":%f, \"power\":%f, \"Motion\":%f\n", Temp, humidity, power, motionCount); //place in string
23     //data sent after stored in string
24     wsoc.send(str);
25     __enable_irq();
26 }
27
28 /*
29 float cosine_func(float x){
30
31     x = 1- 0.5*x*x+ (1/24)*x*x*x*x - (1/720)*x*x*x*x*x*x;
32     return x;
33 }
34 */
35
36 void powerconv()
37 {
38     /* function uses previously collected power data to calibrate the sensor correctly
39     first working out the voltage, current and phase and then real power
40     using known power equations */
41
42     PVoltage = 0.3535*pVoltage - 2.6009;
43     PCurrent = 0.0022*pCurrent - 0.0065;

```

```

44     PPhase = (PI/1024)*pPhase;
45     power = PVoltage*PCurrent*cos(PPhase);
46 }
47
48
49 void lightsignal()
50 {
51     /* function to turn on and off the lights, using the JSON string parsing to
52        find which light has been selected to turn off */
53
54     int lightid = (int)v.get("light_ID").get<double>();
55     //string parsing to get light id
56     swapaddress(MR3address);
57     //select correct xbee address
58
59     /*switch function selects correct light and then using the "status"
60        variable from the JSON string sets light to on or off */
61     switch( lightid ) {
62         case 1:
63             if(v.get("status").get<double>() == 1) {
64                 //call xbee and turn light on
65                 myxbee.SENDLIGHTONREQMOD();
66
67             } else {
68                 //call xbee and turn light off
69                 myxbee.SENDLIGHTOFFREQMOD();
70             }
71             break;
72         case 2:
73             if(v.get("status").get<double>() == 1) {
74                 //call xbee and turn light on
75             } else {
76                 //call xbee and turn light off
77             }
78             break;
79         default:
80             break;
81     }
82 }
83
84
85 void thermoupdate()
86 {
87     /* updates thermostat setting to correct variable from website */
88     thermostat = v.get("setting").get<double>();
89 }
90
91
92 void recievedata()
93 {
94     /* Function that is called when data is recieved from the websocket,
95        the string is then copied to a pointer type and parsed (using
96        JSON functions) to determine the type and then calls other
97        functions depending on the type
98     */
99
100     char * json = (char*) malloc(strlen(recv)+1);
101     strcpy(json, recv);
102
103     string err = picojson::parse(v, json, json + strlen(json));
104     //error function
105
106     /*move the 'type' in JSON string into a string type which is then compared to
107        expected types from the server, after which actions are taken accordingly */
108     char type[200];
109     strcpy(type,v.get("type").get<string>().c_str());
110     printf("%s",type);
111     if( strcmp(type, "light_request") == 0) {
112         //request whether light is on or off
113         sprintf(str, "{\"type\":\"light_request_response\" , \"light_ID\":1 , \"status\": %d }", lightStatus);
114         wsoc.send(str);
115     } else if( strcmp(type, "light_control") == 0) {
116         //turn on or off light
117         lightsignal();
118     } else if( strcmp(type, "thermostat_request") == 0) {
119         //request temperature setting of thermostat, which is returned, along with house temperature
120         sprintf(str, "{\"type\":\"thermostat_request_response\" , \"temperature\":%f , \"setting\": %f }", Temp,
121             thermostat);
122         wsoc.send(str);
123     } else if( strcmp(type, "thermostat_control") == 0) {
124         //thermostat setting has been updated from website
125         thermoupdate();
126     }
127
128     return;
129 }

```

code/homeauto-functions.cpp

XBee

API Framework

Figure 8.1: XBee and XBee PRO wireless specifications

Specification	XBee ZNet 2.5	XBee PRO ZNet 2.5
Performance		
Indoor/Urban Range	up to 133 ft. (40 m)	up to 300 ft. (100 m)
Outdoor RF line-of-sight Range	up to 400 ft. (120 m)	up to 1 mile (1.6 km)
Transmit Power Output	2mW (+3dBm), boost mode enabled 1.25mW (+1dBm), boost mode disabled	50mW (+17dBm) 10mW (+10 dBm) for International variant
RF Data Rate	250,000 bps	250,000 bps
Serial Interface Data Rate (software selectable)	1200 - 230400 bps (non-standard baud rates also supported)	1200 - 230400 bps (non-standard baud rates also supported)
Receiver Sensitivity	-96 dBm, boost mode enabled -95 dBm, boost mode disabled	-102 dBm

The API framework consisted of 3 main sections. First is a function specific frame creator which allows an explicitly specified frame to be transmitted to the UART on the XBee. These frame creator functions are able to calculate their own checksums and account for different addresses as well. This allowed the specific frame requests to be transmitted at request. Next a function called `readreply()` was created in order to capture the reply from the XBees and store it in an array for that function. In most cases the reply came almost immediately, thus there was no need to wait at all for the reply. There are different variants of `readreply` which show the reply bytes if needed. Last is the `readpacket()` function which decoded the reply depending on what it is a reply to. This was achieved by scanning the incoming data and picking out a packet from the reply then looking at the frame specifier type which allowed us to see what kind of response it was. Once this had been done, it was possible to then decode the data from those packets in each case. As mentioned earlier the analogue reading which is a remote AT IS command response has 2 frame ids returned depending on which function created it. Sensor Interfaces.

Our sensors use a variety of these features.

The Motion sensors use the analogue inputs. The API frame for measuring the value on the ADC inputs is sent in the `SENDSSENSORREQ()` function. This function also reads the reply and decodes the appropriate bits to get the analogue value back. The ADC input only goes from 0 to 1.2V though so the levels it can read are different from the motion sensor output which is $VDD-0.5$ where VDD is 3.3V. Thus we used a potential divider to scale the voltage down around 0.8V for motion and 0.5V for no motion. This allows us to see when motion has been detected.

Light

In order to control the light switches with the XBee, we had to implement extra digital logic. This was controlled from two sides, where the physical light switch was one and the digital output from the XBee was the other. In order to service the web socket requests the digital IO functions for the XBee were called by using either `SENDLIGHTONREQ()` or `SENDLIGHTOFFREQ()`. These functions either pulled the pin high or low depending on which one was called. This gave the other input to the XOR gate which allows the triac circuit to turn the light on or off accordingly. The other job the XBee has is to measure the current output of the XOR gate in order to get the current status of the light. This is needed in order to update the graphic on the website relating to the relevant light toggle switch. In the end two ways were attempted to accomplish this.

The first was a polling based method which we ended up using; this method uses a similar analogue function as the motion sensor readings used. This was done in order to save a bit of time when writing the code, as the speed of detection was not an issue. When the function `SENDLIGHTSTATUS()` is called the XBee calls for an analogue reading which is a scaled down version of the digital level on the XOR gate. As well as giving an analogue reading a different frame ID is returned in order to differentiate between motion sensor readings and light switch readings. These are then loaded into global variables and passed via web socket. Our second method was to use UART interrupts along with the XBee pin change detection feature. The XBee can detect rising and falling edges on its monitored digital inputs. When an edge occurs the corresponding API frame is transmitted. The interrupts on the UART allow the MBED to listen for this reply which could come in at any time. Unfortunately the result turned out to be buggy as the modified serial library with a circular buffer had strange behaviour. `getc()` would not read the characters in the correct order, whilst using a function called `rxGetLastChar()` the operation worked fine until the buffer filled up. After flushing this buffer the UART only read in 5 or 6 bytes at a time. The standard serial library worked fine with the interrupts on the UART however not all of the XBee frames can be captured correctly

as there is only a 16 byte buffer on the FIFO. Thus this idea was not used in the end. Since we chose to go down the polling route, there is a 3 second lag on how long it takes for the graphic on the webpage to update.

Power

The power sensor interfaced slightly differently to the other sensors. As there is a ATmega microcontroller on the sensor circuit it was decided that transferring the data as 16 bit short integers over UART would be most efficient. This was the case as the ADCs in the XBee are only 10 bit and would introduce more error into the measurements. Due to the serial interrupts not working as expected a workaround was devised, as it was not possible to just send a raw UART frame we leveraged the ATmega controller to help us. The power readings are taken on a regular poll using a standard timer in the MBED like every other function. When this timer activates the digital output on the power XBee is pulled high. This causes an interrupt to occur on the rising edge of the signal which causes the ATmega controller to explicitly send a new API frame containing raw data with the numbers to the co-ordinator XBee. A small wait happens in between these two functions in order to make sure that the circular buffer is filled with the UART data. After receiving and decoding the data, a request to set the XBee digital pin low is sent, and the sensor will be ready for another reading the next time `SENDPOWERREQ()` is called.

Issues with interrupts

At the beginning the ideal vision of how we would have handled XBee communication would have been to use ISRs to service requests as needed. For simplicities sake due to timing we decided that a polling based method would be more robust and just as viable. In order to use an interrupt based approach we would have to have had full interrupt capability on all of the UART ports so that ISRs could be activated upon the required JSON strings entering the system. Taking a poll based approach we used the timers on the Cortex-M3 in order to activate interrupts at specified time intervals. One problem with this was the fact that whilst one ISR was running, it could not be interrupted at all, so nested interrupts were not possible with this approach. If the ISR was interrupted again whilst sending out an API frame to the UART it would simply be discarded as invalid data. If an interrupt occurred when the XBee received a frame, that data would then have been missed as well and the whole thing would have had to have been done again. Thus all interrupts were disabled globally for the duration of the XBee ISRs. This was done with the function calls `__disable_irq()` and `__enable_irq`. This ultimately limits the amount of processing that can occur on the CPU core as the ISRs are effectively blocking functions, and thus nothing else can run whilst they do. In the case of the power routine, there is a 200ms wait in there which squanders a lot of cycles however it has to be there in order to capture the response, thus in this instance an interruptable UART would have been superior. With polling comes inevitable latency between polls as well which can be several seconds. This is most noticeable when polling for the light context at which point the true context isn't updated for up to 3 seconds if the most recent poll has just passed. This causes a slight glitchyness on the web interface, however after the next poll it would correct itself.

Code

```
1 #include "mbed.h"
2 #include <vector> // alternate way of getting array length as standard wasnt working.
3 #include "MODSERIAL.h"
4 #include "xbee.h"
5
6
7 #define ATIDlength 50
8 #define ATISLENGTH 18
9 #define ATresponlength 9
10 #define MAXMESSAGELENGTH 100
11
12 DigitalOut myled(LED1);
13
14 extern MODSERIAL pc;
15 extern MODSERIAL xbee;
16
17
18 char ATIDplaceholder[ATIDlength];
19 //char BOSSaddress[8]= {0x00,0x13,0xA2,0x00,0x40,0x8B,0x71,0xDE};
20 //char MRladdress[8]= {0x00,0x13,0xA2,0x00,0x40,0x8B,0x71,0x55};
21 char * MRladdress;
22 char * test;
23 vector <char> comvec;
24 int hextoint(char MSB, char LSB)
25 {
26
27     int total=0;
28
29     total=(int)MSB*256+(int)LSB;
30
31     return total;
32 }
33
34
35 void XBEE::swapaddress(char address[]){
36 //address pointed to becomes the one passed in argument.
37 MRladdress=&address[0];
38
39 }
40
```

```

41 char * XBEE::findcurrentaddress(){
42
43 return MRladdress;
44
45 }
46
47
48 char XBEE::checksum(char hexchar[])
49 { //calculate checksum.
50   int length;
51   int calculatedchecksum;
52   int total=0;
53
54   length = (int)hexchar[1]+(int)hexchar[2];
55
56   for(int i=3; i<3+length; i++) { //ignore delimiter and length bytes.
57     total=total+(int)hexchar[i];
58 //pc.printf("%2x\n",hexchar[i]);
59   }
60   //pc.printf("%d\n",total);
61
62   calculatedchecksum=0xFF-(total & 0xFF);
63   //pc.printf("%2x\n",calculatedchecksum);
64
65
66   return calculatedchecksum;
67 }
68
69
70 int XBEE::lastpacketpos(char packet[], int arraylength)//find last packet position in big reply string
71 { //improvement -- detect all packet positions not just last one.
72
73   int currentlastpos=0;
74
75   //pc.printf("array length is...%d\n",arraylength);
76
77   for(int i=0; i<arraylength; i++) { //static array, so valid method.
78
79     if(packet[i]==0x7E)
80       currentlastpos=i;
81   }
82
83   return currentlastpos;
84 }
85
86
87 void XBEE::readpacket(char packetstring[],int packetpos,int packetstringlength)
88 {
89
90   if(packetstringlength != 0) {
91     char snip[128]; //makes a copy of current packet, a preparation for detecting multiple packets in returned
92     string.
93     enum packettype {ATcommandresponse,LocalATcommand,Modemstatus,Txreply,RemoteATcommandresponse,RemoteTxreply};
94     enum packettype packetstatus;
95     enum ATcommand {ATIS,ATID}; //supported AT commands
96     enum ATcommand command;
97     int bytes =packetstring[packetpos+1]*256+packetstring[packetpos+2];
98     //pc.printf("number of bytes in packet is %d\n",bytes);
99     char status = packetstring[packetpos+3]; //get frame specifier
100    pc.printf("status is %2x\n",status);
101    for(int j=0; j<bytes+2+1; j++) { //make a copy of current packet
102      snip[j]=packetstring[packetpos+j];
103    }
104
105    char calculatedchecksum=checksum(snip);
106    //pc.printf("calculated checksum is %2x\n",calculatedchecksum);
107    //pc.printf("given checksum is %2x\n",packetstring[packetpos+bytes+2+1]);
108    bool valid;
109
110    if(calculatedchecksum==packetstring[packetpos+bytes+2+1]) {
111      switch (status) { //assign meaningful specifiers
112
113        case 0x88:
114          packetstatus=ATcommandresponse;
115
116          pc.printf("AT command response packet detected\n");
117          break;
118        case 0x8A:
119          packetstatus=Modemstatus;
120          break;
121        case 0x08:
122          packetstatus=LocalATcommand;
123          pc.printf("LOCAL AT command response packet detected\n");
124          break;
125        case 0x8B:
126          packetstatus=Txreply;
127          pc.printf("TX response packet detected\n");
128          if(packetstring[packetpos+2+5]==0x00)
129            pc.printf("packet was sent sucessfully\n");
130          else
131            pc.printf("Something went wrong\n");
132          break;

```

```

133         case 0x97:
134             packetstatus=RemoteATcommandresponse;
135
136             pc.printf("AT command equivalent is %c%c\n",packetstring[packetpos+2+2+8+2+1],packetstring[
137 packetpos+2+2+8+2+2]); //2/2/8/2===length offset/cmd type and frameID/64bit address/16 bit address
138             char pt1,pt2;
139             pt1=packetstring[packetpos+2+2+8+2+1];
140             pt2=packetstring[packetpos+2+2+8+2+2];
141             if(pt1=='I' && pt2=='S') {
142                 command=ATIS;
143             }
144
145             break;
146         case 0x90:
147             int messagestartpos =0;
148             int datapos=0;
149             packetstatus=RemoteTxreply;
150             pc.printf("REMOTE TX PACKET DETECTED");
151             messagestartpos=packetpos+2+2+8+2+1;
152             pc.printf("char %c%c%c \n", packetstring[messagestartpos],packetstring[messagestartpos+1],
153 packetstring[messagestartpos+2]);
154             if(packetstring[messagestartpos]=='p' && packetstring[messagestartpos+1]=='w' && packetstring[
155 messagestartpos+2]=='r') { //condition for power and not just a standard message
156
157                 datapos=messagestartpos+3;
158                 pc.printf("VCP is %2x%2x%2x%2x%2x\n",packetstring[datapos],packetstring[datapos+1],
159 packetstring[datapos+2],packetstring[datapos+3],packetstring[datapos+4],packetstring[datapos+5]);
160                 //get raw values from power uC.
161                 pVoltage=(int)packetstring[datapos]*256+(int)packetstring[datapos+1];
162                 pCurrent=(int)packetstring[datapos+2]*256+(int)packetstring[datapos+3];
163                 pPhase=(int)packetstring[datapos+4]*256+(int)packetstring[datapos+5];
164
165                 } else {
166                     //IMPROVEMENT -- implment message printer
167                     pc.printf("standard message\n");
168                 }
169
170             }
171
172             if(command==ATIS) {
173                 pc.printf("SENSOR READING REQUESTED\n");
174                 pc.printf("%2x%2x\n",packetstring[packetpos+2+bytes-1],packetstring[packetpos+2+bytes]);
175                 int num=(packetstring[packetpos+2+bytes-1]*256)+packetstring[packetpos+2+bytes];
176                 float voltage=0;
177                 voltage=(float)num/1023*1.2;
178                 pc.printf("rawis %d\n",num);
179                 pc.printf("voltage is %fV\n",voltage);
180                 if(voltage<0.72&&packetstring[packetpos+2+2]==0x01/*gives us Frame ID*/){
181                     pc.printf("movement detected\n");
182                     motionCount++;
183                 }
184                 else if(voltage >=0.72 && packetstring[packetpos+2+2]==0x01){ //no motion detected if voltage is 0.8V
185                     pc.printf("NADA\n");
186                 }
187                 else if(voltage < 0.5 && packetstring[packetpos+2+2]==0x02){
188                     pc.printf("Light is off\n"); //if frame ID is 2 then it must have been a light packet.
189                     lightStatus=false;
190                 }
191
192                 }
193                 else if(voltage>= 0.5 && packetstring[packetpos+2+2]==0x02){
194                     pc.printf("Light is on \n");
195                     lightStatus=true;
196                 }
197
198             }
199
200             else {
201                 pc.printf("the packet is messed up in some way\n");
202             }
203
204         }
205     }
206 }
207 }
208
209 char * XBEE::createmessagepacket(char message[],char address[])
210 {
211
212     static char packet[MAXMESSAGELENGTH +10];
213     short totallength=0;
214     short length = 0;
215     short setuplength = 17; //number of bytes required for address+options etc.
216     while(message[length]!='\0') {
217         length++;
218     }
219
220     pc.printf("length of input message data is %hd\n",length);
221

```

```

222     totallength=length+setuplength;
223     pc.printf("number of bytes in this packet is%hd\n",totallength);
224
225     packet[0]=0x7E;
226     packet[1]=(0xFF00 & (totallength-3)); //-3 for frame delimiters and len MSB/LSB
227     packet[2]=(0x00FF & (totallength-3));
228     packet[3]=0x10;
229     packet[4]=0x01;
230     packet[5]=address[0];
231     packet[6]=address[1];
232     packet[7]=address[2];
233     packet[8]=address[3];
234     packet[9]=address[4];
235     packet[10]=address[5];
236     packet[11]=address[6];
237     packet[12]=address[7];
238     packet[13]=0xFF;
239     packet[14]=0xFE;
240     packet[15]=0x00;
241     packet[16]=0x00;
242     for(int i=0; i<length; i++) {
243         packet[setuplength+i]=message[i];
244         pc.printf("%c was added to message\n",packet[setuplength+i]);
245     }
246     packet[setuplength+length]=checksum(packet);
247
248     for(int i=0; i<totallength+1; i++) { //+1 to get chesum output.
249
250         pc.printf("created packet byte %d is %2x\n",i+1,packet[i]);
251
252     }
253
254     return packet;
255 }
256
257 void XBEE::readreply(char reply[], int &replylength)//if/for implementation, doesnt work well unless reading something
258     very specific.
259 {
260     replylength =0;
261     if(xbee.readable()) {
262
263         for(int i=0; i<18; i++) {
264             reply[i]=xbee.getc();
265             //pc.printf("reply byte %d iz %2x\n",replylength,reply[i]);
266
267             replylength++;
268         }
269
270     }
271 }
272
273 void XBEE::readreply2(char reply[], int &replylength)//shows bytes read in.
274 {
275     replylength =0;
276     while(xbee.readable()) {
277
278
279         reply[replylength]=xbee.getc();
280         pc.printf("reply byte %d iz %2x\n",replylength,reply[replylength]);
281
282         replylength++;
283
284     }
285
286
287
288
289     pc.printf("Length of reply is %d\n",replylength);
290 }
291
292
293
294 void XBEE::readreply3(char reply[], int &replylength)//no bytes shown, so completes a lot quicker as printf is bad for
295     interrupts.
296 {
297     replylength =0;
298     while(xbee.readable()) {
299
300
301         reply[replylength]=xbee.getc();
302
303         replylength++;
304
305     }
306
307
308
309
310
311 }
312

```

```

313
314
315
316 char * XBEE::createsensorpacket(char address[])
317 {
318
319     static char packet[ATISLENGTH];
320     short totallength=18;//number of bytes required for address+options etc.
321
322
323
324     //pc.printf("number of bytes in this packet is%d\n",totallength);
325
326     packet[0]=0x7E;
327     packet[1]=(0xFF00 & (totallength-3));//-3 for frame delimiters and len MSB/LSB
328     packet[2]=(0x00FF & (totallength-3));
329     packet[3]=0x17;
330     packet[4]=0x01;
331     packet[5]=address[0];
332     packet[6]=address[1];
333     packet[7]=address[2];
334     packet[8]=address[3];
335     packet[9]=address[4];
336     packet[10]=address[5];
337     packet[11]=address[6];
338     packet[12]=address[7];
339     packet[13]=0xFF;
340     packet[14]=0xFE;
341     packet[15]=0x00;
342     packet[16]=0x49;//I
343     packet[17]=0x53;//S
344     packet[totallength]=checksum(packet);
345
346
347
348     return packet;
349
350 }
351
352
353 char * XBEE::createlightStatuspacket(char address[]) //could be altered to void if needed
354 {
355
356
357     static char packet[ATISLENGTH];
358     short totallength=18;//number of bytes required for address+options etc.
359
360
361
362     //pc.printf("number of bytes in this packet is%d\n",totallength);
363
364     packet[0]=0x7E;
365     packet[1]=(0xFF00 & (totallength-3));//-3 for frame delimiters and len MSB/LSB
366     packet[2]=(0x00FF & (totallength-3));
367     packet[3]=0x17;
368     packet[4]=0x02;
369     packet[5]=address[0];
370     packet[6]=address[1];
371     packet[7]=address[2];
372     packet[8]=address[3];
373     packet[9]=address[4];
374     packet[10]=address[5];
375     packet[11]=address[6];
376     packet[12]=address[7];
377     packet[13]=0xFF;
378     packet[14]=0xFE;
379     packet[15]=0x00;
380     packet[16]=0x49;//I
381     packet[17]=0x53;//S
382     packet[totallength]=checksum(packet);
383
384
385
386     return packet;
387
388 }
389
390
391 char * XBEE::createlightonpacket(char address[]) //could be altered to void if needed
392 {
393
394     static char packet[ATISLENGTH];
395     short totallength=19;//number of bytes required for address+options etc.
396
397
398
399     //pc.printf("number of bytes in this packet is%d\n",totallength);
400
401     packet[0]=0x7E;
402     packet[1]=(0xFF00 & (totallength-3));//-3 for frame delimiters and len MSB/LSB
403     packet[2]=(0x00FF & (totallength-3));
404     packet[3]=0x17;
405     packet[4]=0x01;

```

```

406     packet[5]=address[0];
407     packet[6]=address[1];
408     packet[7]=address[2];
409     packet[8]=address[3];
410     packet[9]=address[4];
411     packet[10]=address[5];
412     packet[11]=address[6];
413     packet[12]=address[7];
414     packet[13]=0xFF;
415     packet[14]=0xFE;
416     packet[15]=0x02;
417     packet[16]=0x44;
418     packet[17]=0x30;
419     packet[18]=0x05;//make output high
420     packet[totallength]=checksum(packet);
421
422
423
424     return packet;
425 }
426
427 char * XBEE::createlightoffpacket(char address[]) //could be altered to void if needed
428 {
429     static char packet[ATISLENGTH];
430     short totallength=19;//number of bytes required for address+options etc.
431
432     // pc.printf("number of bytes in this packet is%hd\n",totallength);
433
434     packet[0]=0x7E;
435     packet[1]=(0xFF00 & (totallength-3));//-3 for frame delimiters and len MSB/LSB
436     packet[2]=(0x00FF & (totallength-3));
437     packet[3]=0x17;
438     packet[4]=0x01;
439     packet[5]=address[0];
440     packet[6]=address[1];
441     packet[7]=address[2];
442     packet[8]=address[3];
443     packet[9]=address[4];
444     packet[10]=address[5];
445     packet[11]=address[6];
446     packet[12]=address[7];
447     packet[13]=0xFF;
448     packet[14]=0xFE;
449     packet[15]=0x02;
450     packet[16]=0x44;
451     packet[17]=0x30;
452     packet[18]=0x04;//make output low
453     packet[totallength]=checksum(packet);
454
455     return packet;
456 }
457
458 void XBEE::readaddress(char address[]){
459
460     pc.printf("address is %2x%2x%2x%2x%2x%2x%2x \n",address[0],address[1],address[2],address[3],address[4],address[5],
461         address[6],address[7]);
462
463     }//printf current address pointer
464
465 void XBEE::SENDMESSAGE()
466 {
467     int length=0;
468     int replylength=0;
469     char cstr[MAXMESSAGELENGTH];
470     char reply[MAXMESSAGELENGTH];
471     char * msgptr;
472     pc.printf("entermessage\n");
473     pc scanf("%s",cstr);
474     while(cstr[length]!='\0') { //find length of input string
475         length++;
476     }
477
478     pc.printf("length of input is %d\n",length);
479
480     length=length+18;
481
482     msgptr=createmessagepacket(cstr,MRIaddress);
483     pc.printf("message packet created\n");
484
485     if(xbee.writable()) { //send packet down UART

```

```

498         for(int i=0; i<length; i++) {
499             xbee.putc(*(msgptr+i));
500             pc.printf("contents of message packet byte %d is %2x\n", i,*(msgptr+i));
501         }
502     }
503     readreply2(reply,replylength);
504     readpacket(reply,lastpacketpos(reply,replylength),replylength);
505
506     for(int i=0; i<replylength; i++) {
507         pc.printf("reply is %2x and %d\n",reply[i],replylength);
508     }
509 }
510
511 void XBEE::SENDSSENSORREQ()
512 {
513     __disable_irq();
514     int length=19;
515     int replylength=0;
516     char reply[MAXMESSAGELENGTH];
517     char * msgptr;
518     char * temp=findcurrentaddress();
519
520     msgptr=createsensorpacket(MR1address);
521     pc.printf("SENSOR packet created\n");
522
523     if(xbee.writeable()) {
524         for(int i=0; i<length; i++) {
525             //pc.printf("contents of message packet byte %d is %2x\n", i,*(msgptr+i));
526             xbee.putc(*(msgptr+i));
527         }
528     }
529     readreply2(reply,replylength);
530     readpacket(reply,lastpacketpos(reply,replylength),replylength);
531
532     swapaddress(temp);
533     __enable_irq();
534 }
535
536 void XBEE::SENDMOTIONREQ(){
537     __disable_irq();
538     char * temp=findcurrentaddress();
539     swapaddress(MR0address);
540     SENDSENSORREQ();
541     swapaddress(temp);
542     __enable_irq();
543 }
544 void XBEE::SENDLIGHTONREQ()
545 {
546     readaddress(MR1address);
547     int length=20;
548     int replylength=0;
549     char reply[MAXMESSAGELENGTH];
550     char * msgptr;
551
552     msgptr=createlightonpacket(MR1address);
553     //pc.printf("message packet created\n");
554
555     if(xbee.writeable()) {
556         for(int i=0; i<length; i++) {
557             //pc.printf("contents of message packet byte %d is %2x\n", i,*(msgptr+i));
558             xbee.putc(*(msgptr+i));
559         }
560     }
561 }
562
563 void XBEE::SENDLIGHTOFFREQ()
564 {

```



```

591
592
593 int length=20;
594 int replylength=0;
595 char reply[MAXMESSAGELENGTH];
596 char * msgptr;
597
598
599
600 msgptr=createlightoffpacket(MRladdress);
601 pc.printf("message packet created\n");
602
603
604 if(xbee.writeable()) {
605     for(int i=0; i<length; i++) {
606         //pc.printf("contents of message packet byte %d is %2x\n", i,*(msgptr+i));
607         xbee.putc(*(msgptr+i));
608     }
609 }
610
611 readreply2(reply,replylength);
612 readpacket(reply,lastpacketpos(reply,replylength),replylength);
613
614
615
616
617
618
619 }
620
621
622 void xbeeATIDCHANGE()
623 {
624
625     char PANID[2];
626     short hex;
627
628     char hexchar[10]= {0x7E,0x00,0x06,0x08,0x01,0x49,0x44,0x0E,0xE4,0x79};
629
630     pc.printf("what do you want the PANID to be (in hex)?\n");
631     pc scanf("%x", &hex); //hex input
632     pc.printf("%hd\n",hex);
633     pc.printf("%4x\n",hex);
634     PANID[1]=hex/256; //shift left bvy 8 bits so that can fit in char(0xFF00 is too big).
635     PANID[0]=hex & 0x00FF;
636     hexchar[7]=PANID[1];
637     hexchar[8]=PANID[0];
638     hexchar[9]=checksum(hexchar);
639
640
641
642     pc.printf("BEGIN");
643     for(int i=0; i<10; i++) {
644         pc.printf("%2x\n",hexchar[i]);
645     }
646
647
648
649     pc.printf("END");
650
651     if(xbee.writeable()) {
652         for(int i=0; i<10; i++) {
653             xbee.putc(hexchar[i]);
654         }
655     }
656
657 }
658
659
660
661
662
663 }
664
665
666
667
668
669
670 void XBEE::xbeeATID()
671 {
672
673     char hexchar[8] = {0x7E,0x00,0x04,0x08,0x01,0x49,0x44,0x69};
674     char temp;
675     int length=0;
676     checksum(hexchar);
677     if(xbee.writeable()) {
678         for(int i=0; i<=7; i++) {
679             xbee.putc(hexchar[i]);
680         }
681     }
682
683

```

```

684 }
685
686
687
688 for(int i=0; i<ATIDlength; i++) {
689     if(xbee.readable()) {
690         temp=xbee.getc();
691         ATIDplaceholder[i]=temp;
692         comvec.push_back(temp);
693         pc.printf("%2x\n",temp);
694
695     }
696 }
697
698 }
699
700 pc.printf("done\n");
701 while(ATIDplaceholder[length]!='\0') {
702     length++;
703 }
704 pc.printf("stringlength size is%d\n",length);
705 pc.printf("vector size is .... %d\n",comvec.size());
706 pc.printf("last packet pos starts at ... %d\n",lastpacketpos(ATIDplaceholder,comvec.size()));
707 readpacket(ATIDplaceholder,lastpacketpos(ATIDplaceholder,comvec.size()),comvec.size());
708 comvec.clear();
709
710 }
711
712
713 char * XBEE::createlightonpacketMOD(char address[],requestType request) //could be altered to void if needed
714 {
715
716     static char packet[ATISLENGTH];
717     short totallength=19;//number of bytes required for address+options etc.
718
719
720
721     // pc.printf("number of bytes in this packet is%d\n",totallength);
722
723     packet[0]=0x7E;
724     packet[1]=(0xFF00 & (totallength-3));//-3 for frame delimiters and len MSB/LSB
725     packet[2]=(0x00FF & (totallength-3));
726     packet[3]=0x17;
727     packet[4]=0x00;
728     packet[5]=address[0];
729     packet[6]=address[1];
730     packet[7]=address[2];
731     packet[8]=address[3];
732     packet[9]=address[4];
733     packet[10]=address[5];
734     packet[11]=address[6];
735     packet[12]=address[7];
736     packet[13]=0xFF;
737     packet[14]=0xFE;
738     packet[15]=0x02;
739     packet[16]=0x44;//D
740     packet[17]=0x30;//0
741     packet[18]=0x05;//make output high
742     packet[totallength]=checksum(packet);
743
744     if(request==light){
745         lightStatus=true;
746     }
747     return packet;
748 }
749
750
751 char * XBEE::createlightoffpacketMOD(char address[],requestType request) //could be altered to void if needed
752 {
753
754     static char packet[ATISLENGTH];
755     short totallength=19;//number of bytes required for address+options etc.
756
757
758
759     //pc.printf("number of bytes in this packet is%d\n",totallength);
760
761     packet[0]=0x7E;
762     packet[1]=(0xFF00 & (totallength-3));//-3 for frame delimiters and len MSB/LSB
763     packet[2]=(0x00FF & (totallength-3));
764     packet[3]=0x17;
765     packet[4]=0x00;
766     packet[5]=address[0];
767     packet[6]=address[1];
768     packet[7]=address[2];
769     packet[8]=address[3];
770     packet[9]=address[4];
771     packet[10]=address[5];
772     packet[11]=address[6];
773     packet[12]=address[7];
774     packet[13]=0xFF;
775     packet[14]=0xFE;
776     packet[15]=0x02;

```

```

777     packet[16]=0x44;//D
778     packet[17]=0x30;//0
779     packet[18]=0x04;//make output low
780     packet[totallength]=checksum(packet);
781
782     if(request==light){
783         lightStatus=false;
784     }
785     return packet;
786 }
787 }
788
789 void XBEE::SENDLIGHTONREQMOD()
790 {
791     __disable_irq();
792
793     int length=20;
794     int replylength=0;
795     char reply[MAXMESSAGELENGTH];
796     char * msgptr;
797
798
799
800     msgptr=createlightonpacketMOD(MRladdress,light);
801     pc.printf("light on packet created\n");
802
803
804     if(xbee.writeable()) {
805         for(int i=0; i<length; i++) {
806             //pc.printf("contents of message packet byte %d is %2x\n", i,*(msgptr+i));
807             xbee.putc(*(msgptr+i));
808         }
809     }
810
811     __enable_irq();
812 }
813
814 void XBEE::SENDLIGHTOFFREQMOD()
815 {
816     __disable_irq();
817
818     int length=20;
819     int replylength=0;
820     char reply[MAXMESSAGELENGTH];
821     char * msgptr;
822
823
824
825     msgptr=createlightoffpacketMOD(MRladdress,light);
826     pc.printf("light off packet created\n");
827
828
829     if(xbee.writeable()) {
830         for(int i=0; i<length; i++) {
831             //pc.printf("contents of message packet byte %d is %2x\n", i,*(msgptr+i));
832             xbee.putc(*(msgptr+i));
833         }
834     }
835
836     __enable_irq();
837 }
838
839 void XBEE::SENDPOWERONREQMOD()
840 {
841     __disable_irq();
842
843     int length=20;
844     int replylength=0;
845     char reply[MAXMESSAGELENGTH];
846     char * msgptr;
847
848
849
850     msgptr=createlightonpacketMOD(MRladdress,power);
851     pc.printf("light on packet created\n");
852
853
854     if(xbee.writeable()) {
855         for(int i=0; i<length; i++) {
856             //pc.printf("contents of message packet byte %d is %2x\n", i,*(msgptr+i));

```

```

870         xbee.putc(*(msgptr+i));
871     }
872 }
873
874
875
876
877 __enable_irq();
878
879 }
880
881 void XBEE::SENDPOWEROFFREQMOD()
882 {
883     __disable_irq();
884
885
886     int length=20;
887     int replylength=0;
888     char reply[MAXMESSAGELENGTH];
889     char * msgptr;
890
891
892
893     msgptr=createlightoffpacketMOD(MR1address, power);
894     pc.printf("light off packet created\n");
895
896     if(xbee.writeable()) {
897         for(int i=0; i<length; i++) {
898             //pc.printf("contents of message packet byte %d is %2x\n", i, *(msgptr+i));
899             xbee.putc(*(msgptr+i));
900         }
901     }
902 }
903
904
905
906
907
908
909 __enable_irq();
910 }
911
912 void XBEE::SENDPOWERREQ()
913 {
914     __disable_irq();
915     char * temp=findcurrentaddress();
916     /* for(int i=0; i<8; i++){
917         pc.printf("address is %2X\n", MR1address[i]);
918     }*/
919     swapaddress(MR2address); //point to power XBEE
920     /*for(int i=0; i<8; i++){
921         pc.printf("address iss %2X\n", MR1address[i]);
922     }*/
923     char reply[MAXMESSAGELENGTH];
924     int replylength=0;
925
926     SENDPOWERONREQMOD(); //send a digital on request with no reply for interrupt on rising edge for powrr uC
927
928
929
930     wait(0.2); //wait for serial buffer to fill.
931     //pc.printf("INCOMING MESSAGE\n");
932     readreply3(reply, replylength); //see what the uC replies with when it sends a message back
933     readpacket(reply, lastpacketpos(reply, replylength), replylength); //extract raw data values from API frame.
934
935
936     SENDPOWEROFFREQMOD(); //set the interrupt pin on the uC to low ready for the next call.
937     swapaddress(temp);
938
939     __enable_irq();
940 }
941
942 void XBEE::SENDLIGHTSTATUSREQ() {
943     __disable_irq();
944
945     int length=19;
946     int replylength=0;
947     char reply[MAXMESSAGELENGTH];
948     char * msgptr;
949     char * temp=findcurrentaddress();
950     /*for(int i=0; i<8; i++){
951         pc.printf("address is %2X\n", MR1address[i]);
952     }*/
953     swapaddress(MR3address); //edit for #def power XBEE
954     /*for(int i=0; i<8; i++){
955         pc.printf("address iss %2X\n", MR1address[i]);
956     }*/
957
958
959
960
961     msgptr=createlightStatuspacket(MR1address);
962     pc.printf("LIGHT STATUS packet created\n");

```

```

963
964
965     if(xbee.writeable()) {
966
967         for(int i=0; i<length; i++) {
968             //pc.printf("contents of message packet byte %d is %2x\n", i,*(msgptr+i));
969
970             xbee.putc(*(msgptr+i));
971
972         }
973     }
974     readreply2(reply,replylength);
975     readpacket(reply,lastpacketpos(reply,replylength),replylength);
976
977
978
979     swapaddress(temp);
980     __enable_irq();
981
982 }

```

code/homeauto-xbeeC.cpp

```

1  #ifndef XBEE_H
2  #define XBEE_H
3  #include "mbed.h"
4  #include <vector>
5  #include "MODSERIAL.h"
6
7  //external globals declared in main loop which allow measurements to be passed from this library to the main loop.
8
9  extern int pVoltage;
10 extern int pCurrent;
11 extern int pPhase;
12 extern int motionCount;
13
14 //addresses of the XBees
15
16 extern char MR0address[8]; //motion
17 extern char MR2address[8]; //power
18 extern char MR3address[8]; //light control
19 extern bool lightStatus; //current light status
20
21 class XBEE{
22
23     char checksum(char hexchar[]); //calculates checksum of given packet
24     enum requestType{light,power}; //enumeration for differentiating between light control analogue readings and power readings
25     void readpacket(char packetstring[],int packetpos,int packetstringlength); //decode packet information.
26     char * createmessagepacket(char message[],char address[]); //create API frame for sending ASCII messages
27     void readreply(char reply[], int &replylength);
28     void readreply2(char reply[], int &replylength); //read the reply with all bytes shown
29     void readreply3(char reply[],int &replylength); //read the reply with no prints (good for interrupts).
30     char * createsensorpacket(char address[]);
31     char * createlightStatuspacket(char address[]);
32     char * createlightonpacket(char address[]);
33     char * createlightoffpacket(char address[]);
34     char * createlightoffpacketMOD(char address[],requestType request); //MODIFIED packet so that there is no reply from the receiver XBEE so that it doesn't confusate the ASCII string sent from power XBEE
35     char * createlightonpacketMOD(char address[],requestType request);
36     void readaddress(char address[]);
37     char * findcurrentaddress(); //return pointer to array holding current address.
38
39
40
41
42
43
44
45 public:
46
47     void SENDSENSORREQ(); //normal analogue reading on AD1
48     void SENDMOTIONREQ(); //send packet and decode motion sensor reply
49     void SENDLIGHTONREQ(); //turn light on using D0
50     void SENDLIGHTOFFREQ(); //turn light off using
51     void SENDMESSAGE();
52     void xbeeATID(); //find PAN ID
53     void xbeeATIDCHANGE(); //change PAN ID
54     void SENDLIGHTONREQMOD(); //no reply for power uC
55     void SENDLIGHTOFFREQMOD(); //no reply for power uC
56     void SENDPOWERONREQMOD(); //no reply for power uC
57     void SENDPOWEROFFREQMOD(); //no reply for power uC
58     void SENDPOWERREQ();
59     void SENDLIGHTSTATUSREQ(); //find status of light connected to XBee which is currently pointed to
60     void swapaddress(char address[]);
61     int lastpacketpos(char packet[], int arraylength);
62
63
64
65
66
67

```

```
68 |
69 |};
70 |
71 |#endif
```

code/homeauto-xbee.h

8.1.2 Wireless Power Monitoring

Power Diagram

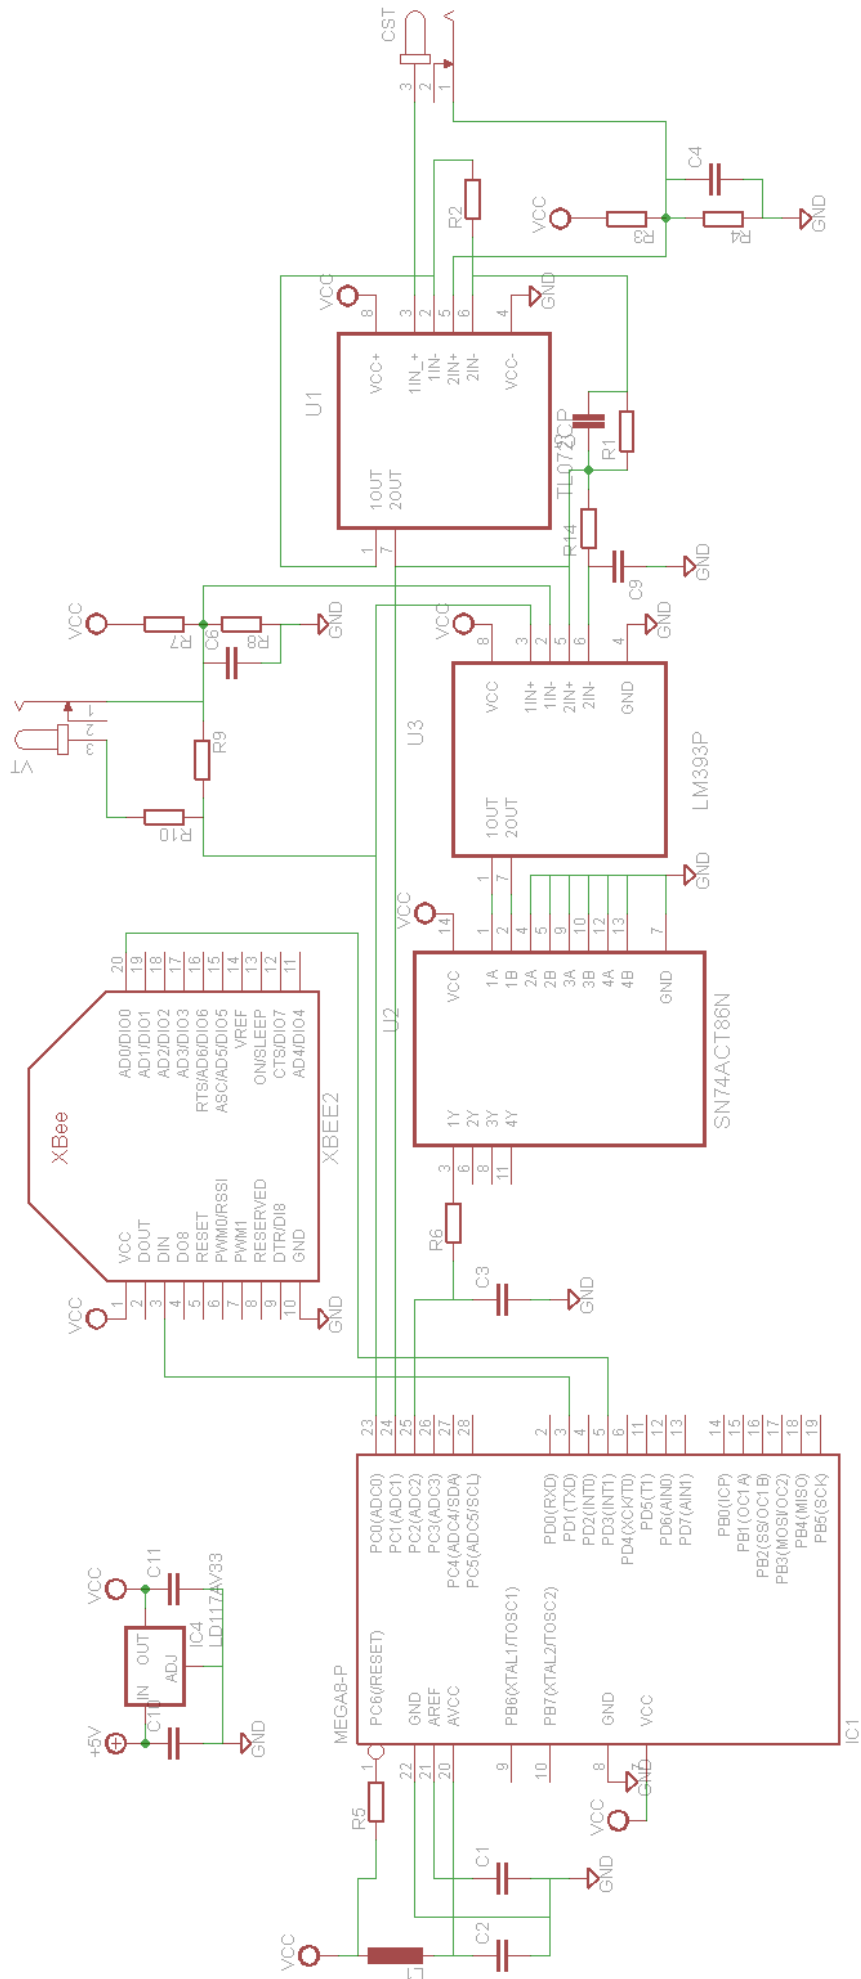


Figure 8.2: Power meter circuit diagram

8.2 Software

8.2.1 Server Code

```
1 #import required libraries
2 import tornado.httpserver
3 import tornado.websocket
4 import tornado.ioloop
5 import tornado.web
6 import time
7 import json
8 import requests
9 import random
10 import functions as mfunc
11 import system_primary_functions as syspf
12
13 #set user count to 0
14 users = 0
15 #custom function to test if a string is JSON
16 def json_test(string):
17     try:
18         json.loads(string)
19         return 1
20     except:
21         return 0
22
23 #handle http
24 class MainHandler(tornado.web.RequestHandler):
25     def get(self):
26         self.write({'type':"message","message":"Hello this is not the correct websocket address. Please add \ws."})
27 #websocket handler
28 class WSHandler(tornado.websocket.WebSocketHandler):
29     #keep track of all users currently connected
30     clients = []
31     def allow_draft76(self):
32         # to use same draft as mbed library
33         return True
34 #on connection open
35 def open(self):
36     #add user to user list
37     self.clients.append(self)
38     print 'new connection'
39     #update user count
40     global users
41     users = users + 1
42     print 'users : %d' % users
43 #send welcome message
44 self.write_message({'type':"message","message":"Hello World"})
45 #log the event
46 with open("tornado_ws.log", "a") as myfile:
47     localtime = time.asctime( time.localtime(time.time()) )
48     myfile.write('%s:new connection : %d users \r\n' % (localtime, users))
49 #on message receipt
50 def on_message(self, message):
51     #log the event
52     print 'message received %s' % message
53     with open("tornado_ws.log", "a") as myfile:
54         localtime = time.asctime( time.localtime(time.time()) )
55         if message != '{"type":"ping"}':
56             myfile.write('%s:message received %s \r\n' % (localtime,message))
57     #test if message is JSON and decide whether to process further
58 if (json_test(message) == 1):
59     print 'message is JSON'
60     #extract data
61     data = json.loads(message)
62     #do different actions based on type of message
63     if data['type'] == 'house_measurement':
64         print "House measurement Data"
65         self.write_message({'type':"message","message":"Sent to Mike GPR"})
66         try:
67             #pass on to rest of clients if temperature data received
68             response = '{"type":"temperature_measurement","temperature":%d}' % (data['Temperature'])
69             print response
70             for client in self.clients:
71                 client.write_message(response)
72         except:
73             print "No temp"
74         try:
75             #send data to the GPR script
76             syspf.get_TS(data['Humidity'])
77         except:
78             print "GPR failed"
79     #post data to database
80     r = requests.post("http://ee-ug1.ee.ic.ac.uk/actual_web2/house_measurement.php", data=data)
81     print r.text
82 #forward all messages with appropriate type
83 elif data['type'] == 'light_request':
84     for client in self.clients:
85         client.write_message(message)
86 elif data['type'] == 'light_request_response':
87     for client in self.clients:
```



```

88         client.write_message(message)
89     elif data['type'] == 'light_control':
90         for client in self.clients:
91             client.write_message(message)
92     elif data['type'] == 'thermostat_control':
93         for client in self.clients:
94             client.write_message(message)
95     elif data['type'] == 'ping':
96         print 'pong'
97         #reply pong to a ping only to the client who sent it
98         response = '{"type": "pong"}'
99         self.write_message(response)
100     elif data['type'] == 'thermostat_set':
101         for client in self.clients:
102             client.write_message(message)
103     elif data['type'] == 'thermostat_request':
104         for client in self.clients:
105             client.write_message(message)
106     elif data['type'] == 'thermostat_request_response':
107         for client in self.clients:
108             client.write_message(message)
109     else:
110         print "Other data"
111     #else:
112     #     print 'message is not JSON'
113
114     def on_close(self):
115         #remove client from list
116         self.clients.remove(self)
117         print 'connection closed'
118         #decrease user count
119         global users
120         users = users - 1
121         print 'users : %d' % users
122         #log it
123         with open("tornado_ws.log", "a") as myfile:
124             localtime = time.asctime( time.localtime(time.time()) )
125             myfile.write('%s:connection closed : %d users \r\n' % (localtime,users))
126
127 application = tornado.web.Application([
128     (r"/", MainHandler),
129     (r'/ws', WSHandler),
130 ])
131
132 #set up to run as process
133 if __name__ == "__main__":
134     http_server = tornado.httpserver.HTTPServer(application)
135     #listen to 8888
136     application.listen(8888)
137     tornado.ioloop.IOLoop.instance().start()

```

code/tornado_ws.py

8.2.2 Timetable Prediction

Choice of day

Testing how many days to take the mode

Obtained from: University of Massachusetts Amherst SMART* dataset

Acc Avg: Is the average accuracy as prediction algorithm is tested on a full 3 months of data

Std.Dev.: is the standard deviation in the average accuracy for each day

len: is the number of days of which the mode is taken to predict the future

s_p: number of past samples used to take the hamming distance of - for example

144 samples corresponds to 12 hours when the sampling is 5 minutes (the case in this data)

s_f: number of future samples taken to then check the accuracy against the know future data

Acc Avg:79.2% Std.Dev.:20% w/day len=2 s_p=144 s_f=144

Acc Avg:79.9% Std.Dev.:22% w/day len=3 s_p=144 s_f=144

Acc Avg:81.5% Std.Dev.:21% w/day len=4 s_p=144 s_f=144

Acc Avg:80.6% Std.Dev.:21% w/day len=5 s_p=144 s_f=144

Acc Avg:80.5% Std.Dev.:21% w/day len=6 s_p=144 s_f=144

Acc Avg:79.5% Std.Dev.:21% w/day len=7 s_p=144 s_f=144

Acc Avg:79.0% Std.Dev.:22% w/day len=8 s_p=144 s_f=144

Acc Avg:79.0% Std.Dev.:22% w/day len=9 s_p=144 s_f=144

Acc Avg:79.0% Std.Dev.:22% w/day len=10 s_p=144 s_f=144

Acc Avg:79.3% Std.Dev.:23% w/day len=11 s_p=144 s_f=144

Acc Avg:79.3% Std.Dev.:23% w/day len=12 s_p=144 s_f=144

Acc Avg:79.3% Std.Dev.:23% w/day len=13 s_p=144 s_f=144


```

Acc Avg:77.1% Std.Dev.:30% w/day len=14 s_p=144 s_f=36
Acc Avg:77.1% Std.Dev.:30% w/day len=15 s_p=144 s_f=36
Acc Avg:77.1% Std.Dev.:30% w/day len=16 s_p=144 s_f=36
Acc Avg:77.1% Std.Dev.:30% w/day len=17 s_p=144 s_f=36
Acc Avg:77.1% Std.Dev.:30% w/day len=18 s_p=144 s_f=36
Acc Avg:77.1% Std.Dev.:30% w/day len=19 s_p=144 s_f=36
Acc Avg:77.1% Std.Dev.:30% w/day len=20 s_p=144 s_f=36
Acc Avg:77.1% Std.Dev.:30% w/day len=21 s_p=144 s_f=36
Acc Avg:77.1% Std.Dev.:30% w/day len=22 s_p=144 s_f=36
Acc Avg:77.1% Std.Dev.:30% w/day len=23 s_p=144 s_f=36
Acc Avg:77.1% Std.Dev.:30% w/day len=24 s_p=144 s_f=36
Acc Avg:77.1% Std.Dev.:30% w/day len=25 s_p=144 s_f=36
Acc Avg:77.1% Std.Dev.:30% w/day len=26 s_p=144 s_f=36
Acc Avg:77.1% Std.Dev.:30% w/day len=27 s_p=144 s_f=36
Acc Avg:77.1% Std.Dev.:30% w/day len=28 s_p=144 s_f=36
Acc Avg:77.1% Std.Dev.:30% w/day len=29 s_p=144 s_f=36
Acc Avg:77.1% Std.Dev.:30% w/day len=30 s_p=144 s_f=36

```

Matlab Timetable Prediction Test Script

```

1 clc;
2 clear all;
3 clear;
4 close all;
5 %load data
6 % load house_data
7 %{
8 load housestate.mat
9 clear t_mod
10 clear total_firings
11 clear p_state
12 %}
13 %%{
14 %setup variables
15 load usualday
16 t_range = dat;
17 house_state = num';
18 clear num;
19 clear dat;
20 %}
21 %load threemonths
22 % set all constants
23 day_samples = 96;%288
24 ratio = 288/day_samples; %ratio to 288 - used if non-5minute intervals
25 sample_p = 144/ratio; %half a day
26 sample_f = 144/ratio; %half a day
27
28 %% day prediction
29
30 %set time offset
31 time_o = 90/ratio; %equivalent to 7:30 AM
32
33 %store all data in a 3d array for optimisation
34 acc_3d = [1 1 1];
35
36 day_to_compare = 4; % days to perform a mode on
37 days = 2:length(house_state)/day_samples-1; % days to cycle through
38 for len=day_to_compare
39     for sample_p = round(day_samples*(1:40)/40)
40         for sample_f = round(day_samples*(1:40)/40)%[36 72 144]
41             %set timing variable
42             acc = 0;
43             for day=days
44                 day_index = time_o+(day-1)*day_samples+1; %day with time offset
45                 %store data for future - i.e. to check for accuracy
46                 day_f = house_state(day_index:day_index+sample_f-1);
47                 %store data from past to compute Hamming distances
48                 day_p = house_state(day_index-sample_p:day_index-1);
49                 %prepare data on states - removing excess data
50                 states = house_state;
51                 states(day_index-sample_p:day_index+sample_f-1) = 3; %set to 3 as this is impossible state
52                 %loop through each day to calculate hamming distance of each
53                 %day relative to each other
54                 for i=1:length(states)/day_samples-1
55                     % disp(sprintf('Day: %d',i))
56                     % disp(datestr(t_range(i*day_samples-sample_p+time_o+1)))
57                     % disp(datestr(t_range(i*day_samples+time_o)))
58                     comp = [day_p;states(i*day_samples-sample_p+time_o+1:i*day_samples+time_o)];
59                     D(i) = pdist(comp,'hamming');
60                     % figure;
61                     % subplot(2,1,1), plot(states(i*day_samples-sample_p+time_o+1:i*day_samples+time_o))
62                     % subplot(2,1,2), plot(day_p)
63                     % title(sprintf('Day:%d Hamming distance:%.2f',i,D(i)))

```

```

64         end
65         %get minimum values of D with their indexes
66         [sortedValues,sortIndex] = sort(D,'ascend');
67         maxIndex = sortIndex(1:len); %get minimum hamming distance days
68         %make sure not to go out of range
69         if any(maxIndex == length(states)/day_samples-1)
70             index = find(maxIndex==length(states)/day_samples-1);
71             maxIndex(index) = sortIndex(len+1);
72         end
73         %days_hd stores all the days with the smallest hamming distance
74         days_hd = ones(sample_f,len)';
75         for i=maxIndex
76             %disp(datestr(t_range(i*day_samples+time_o+1)))
77             %disp(datestr(t_range(i*day_samples+time_o+sample_f)))
78             days_hd = [days_hd ;house_state(i*day_samples+time_o+1:i*day_samples+time_o+sample_f)];
79         end
80         %strip first value
81         days_hd = days_hd(2:end,:);
82         days_hd = mode(days_hd');
83         days_hd(days_hd==2) = 1;
84         day_f(day_f==2) = 1;
85         %calc accuracy
86         acc = [acc 100*(1-pdist([days_hd; day_f],'hamming'))];
87         %figure;
88         %subplot(2,1,1), plot(days_hd)
89         %subplot(2,1,2), plot(day_f)
90     end
91     acc = acc(2:end);
92     %print data
93     fprintf('Acc Avg:%2.1f%% Std.Dev.:%2.f%% w/day len=%d s_p=%d s_f=%d\n',mean(acc),std(acc),len,sample_p,
sample_f);
94     acc_3d = [acc_3d;sample_p,sample_f,mean(acc)];
95     end
96 end
97 %store in variable to allow 3d plot
98 acc_3d = acc_3d(2:end,:);

```

general_prediction.m

Python code to predict timetable

```

1  #import all used modules
2  import MySQLdb
3  import gc
4  import numpy as np
5  import scipy as sp
6  from scipy import spatial
7  from scipy import stats
8  from scipy import signal
9  import math
10 import sys
11 import time
12 from datetime import datetime, timedelta
13
14
15 #from past data predicts the expect day states
16 #in: day present states
17 #out: expected future states
18 def get_future_samples(day_p):
19     #initialize distance array
20     D = np.array([])
21     #loop through each day finding the closest days in terms of hamming distance
22     for i in range(1,days-2):
23         #get array to compare
24         comp = np.array([day_p,house_state[i*day_samples-sample_p+time_o:i*day_samples+time_o]][:,:])
25         #store hamming distance in D
26         D = np.append(D,sp.spatial.distance.pdist(comp,'hamming'))
27     #sort by index
28     sortIndex = np.argsort(D)
29     days_hd = []
30     #get days to compare
31     for i in sortIndex[0:day_to_compare]:
32         days_hd.append(house_state[i*day_samples+time_o:i*day_samples+time_o+sample_f])
33     #convert to array
34     days_hd = np.array(days_hd)
35     #get the mode
36     days_hd = sp.stats.mode(days_hd)[0][0]
37     return days_hd
38
39 #custom discrete decimating array function
40 def decimate_array(x, q):
41     x_out = []
42     if not isinstance(q, int):
43         raise TypeError("q must be an integer")
44     if len(x)%q != 0:
45         raise NameError("Cannot decimate!")
46     for i in range(0,len(x)/q):
47         x_out.append(int(sp.stats.mode(x[i*q:(i+1)*q])[0][0]))
48     return x_out

```



```

142 for i,v in enumerate(day_predict):
143     if v==0:
144         day_predict[i] = 2
145     if v==2:
146         day_predict[i] = 0
147 #send to mySQL
148 print "sending to mySQL",
149 #store in correctly named table
150 day_index_name = "Auto-%s" % ((date_object+timedelta(minutes=144*5)).strftime('%A'))
151 print ".",
152 db = MySQLdb.connect(host="127.0.0.1", port=3306, user="root", passwd="London123", db="House_states")
153 cursor = db.cursor()
154 print ".",
155 for i,v in enumerate(day_predict):
156     command = "UPDATE '%s'\nSET state=%d\nWHERE Hour=%d;" % (day_index_name,v,i)
157     cursor.execute(command)
158 print "."
159
160
161 #process other days in the week
162 day_list = range(0,(date_object+timedelta(minutes=144*5)).weekday())
163 day_list.extend(range((date_object+timedelta(minutes=144*5)).weekday()+1,7))
164 #loop through every last 4 days of the week and take the mode of each day
165 for d in day_list:
166     iteration = 0
167     day_states_data = []
168     for i in range(days-1,-1,-1):
169         if datetime.fromtimestamp(int(date[i*day_samples][0])/1000.0).weekday() == d:
170             iteration+=1
171             day_states_data.append(house_state[i*day_samples:(i+1)*day_samples])
172             if iteration >= day_to_compare:
173                 break
174     day_index_name = "Auto-%s" % (datetime.fromtimestamp(int(date[i*day_samples][0])/1000.0).strftime('%A'))
175     day_states_data = decimate_array(np.asarray(sp.stats.mode(day_states_data)[0][0]).astype(int),12)
176     #invert 2 and 0
177     for i,v in enumerate(day_states_data):
178         if v==0:
179             day_states_data[i] = 2
180         if v==2:
181             day_states_data[i] = 0
182     for i,v in enumerate(day_states_data):
183         command = "UPDATE '%s'\nSET state=%d\nWHERE Hour=%d;" % (day_index_name,v,i)
184         cursor.execute(command)
185 #close SQL connection
186 db.close()
187 gc.collect()
188
189
190 print'''
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
'''

```

sql-ttfill.py

8.2.3 Gaussian Process Regression

General Machine Learning (ML) Terminology

Features are the input variables to the algorithm, i.e. if the algorithm is the function that outputs thermostat settings, then the features of the ML algorithm are the inputs or the feature space/domain. Features are inputted as vectors typically denoted by **x**. Our choice of features was dictated by what made logical sense by the data available to us. The features selected are listed below:

1. External Temperature
2. External Humidity
3. Internal Humidity
4. Wind Speed
5. Wind Direction

(Internal temperature was omitted as this would provide unwanted feedback on the feedback from the control system, which the thermostat uses to regulate the internal temperature)

1. **Training data set** - is the set of feature vectors or ‘training examples’ in the feature space that the algorithm uses to fit its regression curve in order to predict the thermostat setting for different conditions.
2. **Target data set** - is the set of outputs (in this context thermostat settings) for each of the training data points.
3. **Model Parameters** - are the set of unknown coefficients of the model that is applied to the data. Linear Regression is an example of a parametric approach in which a model is chosen and then the parameters of the model are optimized to best fit the data.
4. **Hyperparameters** - are important in non- parametric models such as Gaussian Process Regression. Gaussian Process Regression has no direct parameters to find or optimize, but it does have parameters for the prior, i.e. the coefficients of the covariance function. These need to be optimized to fit the data in a similar way to model parameters in parametric regression.
5. **Hypothesis** - also known as the regression curve is the function that is inferred from the training data.

Terms and Definitions for GPR:

1. **Random Variable** – A random variable is a variable that can take on a set of possible values each with an associated probability.
2. **Prior Probability** –the probability distribution that one assigns to a quantity that expresses the uncertainty about the process before any data is observed. In GPR we assume a Gaussian distribution prior.
3. **Likelihood** – the probability of the system parameters given the data that is observed.
4. **Marginal Likelihood** –the likelihood function integrated over certain parameters. In this way the parameters have been ‘marginalized’. It is also known as the ‘model evidence’ because integrating this over the function f gives the probability of the data given the chosen hyperparameters.
5. **Posterior Probability** –The conditional probability that is assigned once relevant data and information is taken into account.
6. **Covariance** - A measure of how much two random variables change with one another.
7. **Stochastic Process** – is a collection of random variables.
8. **Covariance Function** – For a stochastic process or random field, a covariance function describes the spatial covariance of points in the random field located at certain points in space. Covariance functions belong to a family of mapping functions called ‘kernels’.
9. **Multivariate Gaussian distribution** –a generalization of the standard Gaussian or Normal distribution to multiple dimensions. It is described by a mean vector with a separate mean for each dimension, and a covariance matrix (a generalization of variance to multiple dimensions). It gives a similarity measure of each point with every other point in the distribution and is expressed as:

$$\beta \sim N(\mu, \mathbf{K})$$

β is a multi-dimensional random vector, μ is a mean vector defining the mean for each dimension, and \mathbf{K} is the covariance matrix, detailing how much each random variable of \mathbf{x} changes with another, defined as:

$$\mathbf{K} = \begin{bmatrix} \left(\begin{array}{ccc} k(\mathbf{x}_1, \mathbf{x}_1) & \cdots & k(\mathbf{x}_1, \mathbf{x}_n) \\ \vdots & \ddots & \vdots \\ k(\mathbf{x}_n, \mathbf{x}_1) & \cdots & k(\mathbf{x}_n, \mathbf{x}_n) \end{array} \right) \end{bmatrix}$$

The Probability Density Function (PDF) of a Multivariable Gaussian Distribution is:

$$pdf_{\beta} = P(\beta=\mathbf{a}) = (2\pi)^{-\left(\frac{n}{2}\right)} |\mathbf{K}|^{-0.5} \exp\left(-\frac{1}{2}(\mathbf{a} - \mu)^T \mathbf{K}^{-1}(\mathbf{a} - \mu)\right) \quad (8.1)$$

1. **Gaussian Process** - Is a stochastic process whose random variables, associated with each point in the random domain space, have a normal distribution. A Gaussian process can be thought of/visualised as an infinite dimensional Gaussian distribution. It is defined as:

$$f(\mathbf{x}) \sim GP(m(\mathbf{x}), k(\mathbf{x}_i, \mathbf{x}_j))$$

$m(\mathbf{x})$ is the mean function and ' $k(\mathbf{x}_i, \mathbf{x}_j)$ ' is the similarity or covariance function that measures the similarity between two random variables of the Gaussian process. For ease of computation, and without loss of generality we can set $m(x) = 0$ as summing an infinite array of possible functions averages to 0. An important property of a Gaussian Process is that it obeys the consistency requirement; i.e. inspection of a larger set of random variables does not change the distribution of a smaller set. Every finite collection of these random variables can be described with a multivariate Gaussian distribution.

Algorithms and their Complexity:

For this system to operate effectively three algorithms/functions are required; one to search the error curve to find a range of hyperparameter initialisations that find an acceptable minima, a second to train the GPR and a third to run the GPR and return a predicted thermostat setting. In this section ' m ' denotes the number of features and ' n ' the number of examples in the training set. Furthermore $n \gg m$, $n \gg$ any loop counters

Algorithm 1: Hyperparameter Search

Description:

Called every time a significant change in the training data is detected, so fairly infrequently (e.g. the first time the training data reaches its maximum capacity, or when a significant amount of data is changed).

Pseudo code:

1. Import training data into \mathbf{X} and cross validation data into \mathbf{x}^*
2. Define upper and lower bound (ub and lb) as the boundaries within which to initialise θ values. Initialise over some very large range.
3. Grid Search:

while (ub - lb) > 1

increment = (ub - lb) / numb_div

lb_temp = lb

for j = 1->(numb_div - 1)

Initialise θ in range(lb_temp, ub)

Train the GPR

Run the cross validation data through 'GPR Run'

$$\text{AvError}[j] = \frac{1}{n} \sum_{i=1}^n \left\{ \frac{(y_i - \text{ypred}_i) * 100}{y_i} \right\}$$

minIndex = index of minimum element in AvError[j]

ub = lb + minIndex * increment

lb = lb + (minIndex - 1) * increment

1. Write upper bound and lower bound values to SQL database to be used by 'GPR train'

$$[Total\ Complexity \approx \log_{numb_{div}} (ub - lb) * numb_{div} * (O(n) + O(n^2) + O(n^3)) \approx O(n^3)]$$

Algorithm 2: GPR Train

Description:

Called every time the training data set is updated. Calculates the inverse covariance matrix (plus noise) and the optimized hyperparameters, which are then both written back to the SQL database to be used by the 'GPR run' algorithm. As an aside the complexity of the kernel $k(\mathbf{x}_n, \mathbf{x}_n)$ is, considering only multiplications, divisions, subtractions and additions, $(m^2 + m + 1) \approx O(m^2)$

Pseudo code:

1. Import training data from SQL database into \mathbf{X} and the target data into \mathbf{y}
2. Normalize \mathbf{X} : $\forall x_i \text{ in } \mathbf{X}, x_i = \frac{(x_i - \mu)}{\sigma_n}$ (this ensures that all features are equally important and improves performance for gradient descent as each dimension has a similar range of relevant magnitudes)
3. Initialise hyperparameters, θ , by taking values from a uniform distribution with limits found by the hyperparameter search
4. Gradient Descent:

for j = 1->max_iters **or** until $L(\theta)$ not decreasing **or** $L(\theta) < desired_error$

$$\mathbf{C} = \left[\begin{pmatrix} k(\mathbf{x}_1, \mathbf{x}_1) + \sigma_n^2 \delta_{ij} & \cdots & k(\mathbf{x}_1, \mathbf{x}_n) + \sigma_n^2 \delta_{ij} \\ \vdots & \ddots & \vdots \\ k(\mathbf{x}_n, \mathbf{x}_1) + \sigma_n^2 \delta_{ij} & \cdots & k(\mathbf{x}_n, \mathbf{x}_n) + \sigma_n^2 \delta_{ij} \end{pmatrix} \right] [n^2 (2m^2 + 1) \approx O(n^2 m^2)]$$

$$L(\theta) = \frac{1}{2} \mathbf{y}^T \mathbf{C}^{-1} \mathbf{y} + \frac{1}{2} \log(|\mathbf{C}|) + \left(\frac{n}{2}\right) \log(2\pi) [O(n^2) + O(n^3) \approx O(n^3)]$$

for i = 1-> numb.hyperparameters

$$\theta_i = \theta_i - \alpha \frac{\partial L(\theta)}{\partial \theta_i} [n^3 + 2n + 2n^2 \approx O(n^3)]$$

Return θ

1. Using optimized θ values calculate $\mathbf{C} = \left[\begin{pmatrix} k(\mathbf{x}_1, \mathbf{x}_1) + \sigma_n^2 \delta_{ij} & \cdots & k(\mathbf{x}_1, \mathbf{x}_n) + \sigma_n^2 \delta_{ij} \\ \vdots & \ddots & \vdots \\ k(\mathbf{x}_n, \mathbf{x}_1) + \sigma_n^2 \delta_{ij} & \cdots & k(\mathbf{x}_n, \mathbf{x}_n) + \sigma_n^2 \delta_{ij} \end{pmatrix} \right]$
2. Calculate $\mathbf{C}^{-1} [O(n^3)]$
3. Write θ and \mathbf{C}^{-1} to SQL database

$$[Total\ Complexity = 3O(n^3) + O(n^2 m^2) \approx O(n^3)]$$

Algorithm 3: GPR Run

Description:

Algorithm is called every time the embed requests a new thermostat setting. Returns the desired thermostat setting and the confidence with which the prediction is made.

Pseudo code:

1. Import training data into \mathbf{X} , target data into \mathbf{y} , new input feature vector \mathbf{x}^* , the hyperparameters θ and \mathbf{C}^{-1}

$$2. \text{ Calculate } \mathbf{K}_* = \left[\begin{pmatrix} k(\mathbf{x}_*, \mathbf{x}_1) \\ \vdots \\ k(\mathbf{x}_*, \mathbf{x}_n) \end{pmatrix} \right] [n(m^2 + m + 1) \approx O(nm^2)]$$

$$3. \text{ Calculate } \mathbf{K}_{**} = k(\mathbf{x}_*, \mathbf{x}_*) [O(m^2)]$$

$$4. \text{ Predicted Thermostat Setting: } TS = \mathbf{K}_*^T (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{y} [2n^2 \approx O(n^2)]$$

$$5. \text{ Confidence: } TS_{conf} = \mathbf{K}_{**} - \mathbf{K}_*^T (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{K}_* [1 + 2n^2 \approx O(n^2)]$$

$$[Total\ Complexity = O(nm^2) + 2O(n^2) + O(m^2) \approx O(n^2)]$$

Code

```

1  #!/usr/bin/python
2
3  from numpy import *
4  from math import *
5
6
7  #COVARIANCE FUNCTION
8  #Calculates the covariance, i.e. a measure of the similarity between two points
9  #using the error squared or Gaussian kernel
10 def cov_func(x1, x2, M, sig_pow):
11     diff = x1 - x2
12     temp = dot(diff.transpose(), M)
13     k = sig_pow * exp(-0.5 * dot(temp, diff))
14     #print k
15     return k
16
17
18
19
20 #CALCULATE COVARIANCE MATRIX K
21 #Uses the covariance function to calculate the covariance matrix of the multivariate gaussian
22 #distribution of the training data points.
23 def calc_K(X, theta):

```

```

24
25 #Number of Input features
26 m = len(X[0])
27 #Number of data points
28 n = len(X)
29
30 #Initialise matrix M
31 M = zeros((m,m))
32 for i in range(m):
33     M[i][i] = 1/(theta[i+2]*theta[i+2])
34
35 #Initialise then calculate
36 K = zeros((n,n))
37 for i in range(n):
38     for j in range(i, n):
39         K[i][j] = cov_func(X[i].transpose(), X[j].transpose(), M, theta[0])
40         K[j][i] = K[i][j]
41     return K
42
43
44 #CALCULATE 'C'
45 #'C' is just the notation for the covariance matrix K plus a noise term that represents
46 #jitter in the training data
47 def calc_C(K, noise_pow):
48     a = len(K)
49     temp = noise_pow*eye(a)
50     C = K + temp
51     return C
52
53
54 #PARTIAL DIFFERENTIATE C WITH RESPECT TO HYPERPARAMETERS
55 #Function is necessary to carry out gradient descent. Returns a vector containing the
56 #evaluated partial differential of NLML with respect to each of the hyperparameters
57 def grad_C(theta, X, K, C):
58
59     #a is the number of hyperparameters
60     a = len(theta)
61     #n is the number of data points in the training set, note that K and C are nxn
62     n = len(C)
63     #m is the number of data features
64     m = len(X[0])
65
66
67
68 #Initialise gradient of C
69 partial_diff_C = zeros((a,n,n))
70
71 #Find partial diff with respect to signal power
72 partial_diff_C[0] = 2*theta[0]*K
73
74 #Find partial diff with respect to the noise power
75 partial_diff_C[1] = 2*theta[1]*eye(n)
76
77 #Find partial diff with respect to the different scale lengths
78 for k in range(2,a):
79     for i in range(n):
80         for j in range(i, n):
81             diff = (X[i][k-2] - X[j][k-2])**2
82             l = theta[k]
83             temp = (diff/(l**3))
84             partial_diff_C[k][i][j] = temp*K[i][j]
85             partial_diff_C[k][j][i] = partial_diff_C[k][i][j] #Note that C and K are symmetric, therefore
86             only need to calculate half the elements
87
88 return partial_diff_C
89
90
91 #CALCULATE THE GRADIENT OF THE NEGATIVE LOG MARGINAL LIKLIHOOD (NLML) WITH RESPECT TO ONE OF THE FEATURES
92 #Calculates the partial differential of the NLML with respect to a single hyper parameter
93 def calc_pdiff_nlml(C, y, pdiff_C):
94
95     inv_C = linalg.inv(C)
96     diff_L = 0.5*(trace(dot(inv_C, pdiff_C)) - dot(y.transpose(), dot(inv_C, dot(pdiff_C, dot(inv_C,y)))))
97     return diff_L
98
99
100
101 #CALCULATE THE NEGATIVE LOG MARGINAL LIKLIHOOD (NLML)
102 #Calculates the Negative Log Marginal Likelihood
103 def calc_nlml(C, y):
104     n = len(C)
105     L = 0.5*(log(linalg.det(C)) + dot(y.transpose(),dot(linalg.inv(C), y)) + n*log(2*pi))
106     return L
107
108
109 #GRADIENT DESCENT
110 #Carry out gradient descent to find a minima to optimize the hyper parameters with respect to
111 #minimising the NLML.
112 def gradient_descent(alpha, init_theta, X, y, error_level, max_iters):
113     m = len(init_theta)
114     L_hist = ones(max_iters+2)
115     pdiff_L = zeros(m)

```

```

116 theta = init_theta
117 j = 0
118
119 K = calc_K(X, theta)
120 C = calc_C(K, theta[1])
121 L_hist[j] = calc_nllm(C, y)
122
123 successful = True
124 while(L_hist[j] > error_level):
125     j = j+1
126     K = calc_K(X, theta)
127     C = calc_C(K, theta[1])
128     L_hist[j] = calc_nllm(C, y)
129     pdiffCs = grad_C(theta, X, K, C)
130     for k in range(m):
131         p_diffL = calc_pdiff_nllm(C, y, pdiffCs[k])
132         theta[k] = theta[k] - alpha*p_diffL
133
134
135 if(j>max_iters):
136     print("Failure to converge below desired error level after number of iterations =")
137     print j
138     print("Consider increasing value of alpha to decrease convergence time")
139     successful = False
140     break
141
142 if(L_hist[j] > L_hist[j-1]):
143     print("Failure to converge below required error level, nllm not decreasing with each iteration, local minima
reached, consider decreasing alpha")
144     print("Number of iterations =")
145     print j
146     successful = False
147     break
148
149 if(successful == True):
150     print("Convergence below desired error level successful after number of iterations =")
151     print j
152
153 return (theta, successful, L_hist[j], L_hist, j)
154
155
156 #CALCULATE K_*
157 #Calculate the covariance of the input data point with each the points in the training set
158 def calc_K_star(x, X, theta):
159
160     #Number of Input features
161     m = len(X[0])
162     #Number of data points
163     n = len(X)
164     #Initialise matrix M
165     M = zeros((m,m))
166     for i in range(m):
167         M[i][i] = 1/(theta[i+2]*theta[i+2])
168     K_star = zeros(n)
169
170     for j in range(n):
171         K_star[j] = cov_func(x.transpose(), X[j].transpose(), M, theta[0])
172     return K_star
173
174
175 #CALCULATE EXPECTED THERMOSTAT SETTING
176 def calc_E_TS(inv_C, K_star, y):
177     expected_TS = dot(K_star, dot(inv_C, y))
178     return expected_TS
179
180 #NORMALIZE INPUT DATA BEFORE CARRYING OUT GRADIENT DESCENT
181 #Normalisation is necessary because otherwise differences in the order of magnitude between
182 #different features would result in a warped (and much slower) path of descent
183 def normalize_data(X):
184     X = X.transpose()
185     numfeatures = len(X)
186     numdata = len(X[0])
187     norm_X = zeros((numfeatures, numdata))
188
189     for j in range(numfeatures):
190         u = mean(X[j])
191         u = u*ones(numdata)
192         stan_dev = std(X[j])
193         stan_dev
194         norm_X[j] = (1/stan_dev)*(X[j] - u)
195     return norm_X.transpose()
196
197 #NORMALIZE TRAINING AND INPUT DATA ACCORDING TO TRAINING DATA NORMALISATION
198 #Normalizes the training data and also the input data. Note the training data is normalized
199 #first and then the mean and std are taken and used to normalize the input. This ensures the same
200 #transformation occurs on the input data as on the training data and therefore that the relationship
201 #with the thermostat setting is the same.
202 def normalize_input(X, x):
203     X = X.transpose()
204     numfeatures = len(X)
205     numdata = len(X[0])
206     norm_X = zeros((numfeatures, numdata))
207     norm_x = zeros(numfeatures)

```

```

208
209     for j in range(numfeatures):
210         u = mean(X[j])
211         u_vec = u*ones(numdata)
212         stan_dev = std(X[j])
213         stan_dev
214         norm_X[j] = (1/stan_dev)*(X[j] - u_vec)
215     norm_x[j] = (1/stan_dev)*(x[j] - u)
216     return norm_X.transpose(), norm_x
217
218
219 #Function used to find the index and value of an element in an
220 #array that is closest to the argument 'value'. Example application
221 #is finding closest time to the present time in a database.
222 def find_nearest(array,value):
223
224     min_index = 0
225     min_diff = array[min_index] - value
226     diff = zeros(len(array))
227     min_value = array[min_index]
228
229     for j in range(len(array)):
230         diff[j] = abs(array[j] - value)
231         if(diff[j] <= min_diff):
232             min_index = j
233             min_diff = diff[min_index]
234             min_value = array[j]
235
236     return min_index, min_value, min_diff
237
238 #Calculates a 'similarity' metric for two vectors. An example application is in finding entries in the training data
239 #base whose
240 #external conditions best match a new training data point that needs to be swapped in.
241 def calc_similarity(a, b):
242     if(len(a) != len(b)):
243         print("Error: cannot calculate similarity between vectors as do not have the same number of elements")
244         sum = 0
245
246     for j in range(len(a)):
247         sum = sum + abs(a[j] - b[j])
248
249     return sum

```

code/gpr-functions.py

```

1 #!/usr/bin/python
2 import MySQLdb
3 import math
4 import random
5 import sqlite3
6 import time
7 from functions import *
8 from websocket import create_connection
9 from datetime import datetime
10 from scipy import *
11 from numpy import *
12
13
14
15 #GPR_RUN:
16 #Function reads in training data, trained hyperparameters and inverse C and uses them to calculate
17 #a thermostat estimate. This is then returned to be sent via websockets back to the embed module.
18 def active_TS(x):
19     #print("active_TS function entered, calculating TS value using Gaussian Process Regression...")
20     db = MySQLdb.connect("127.0.0.1","mike","ChickenKorma","GPR_Training_Data" )
21     cur = db.cursor()
22     #Read in trained hyperparameters to theta from SQL database
23     cur.execute("SELECT * FROM trained_theta_values")
24     row = cur.fetchall()
25     numb_hp = len(row[0])
26     theta = zeros(numb_hp)
27     for j in range(numb_hp):
28         theta[j] = row[0][j]
29     #Read in inv_C from SQL database
30     cur.execute("SELECT * FROM inv_C")
31     rows = cur.fetchall()
32     numrows = int(math.sqrt(len(rows)))
33     inv_C = ones((numrows, numrows))
34     for j in range(numrows):
35         for i in range(numrows):
36             inv_C[j][i] = rows[i][2]
37     #Import the training data required in calculation
38     cur.execute("SELECT * FROM training_data")
39     rows = cur.fetchall()
40     numrows = len(rows)
41     numcols = len(rows[0])
42     X = zeros((numrows, (numcols-1)))
43     y = zeros((numrows, 1))
44     for j in range((numrows)):
45         y[j] = rows[j][0]
46     for j in range(numrows):
47         for i in range(1, numcols):

```

```

48     X[j][(i-1)] = rows[j][i]
49
50 #Normalize Data
51 X, x = normalize_input(X, x)
52
53 #Calculate K_star and then calculate expected TS value given input x
54 K_star = calc_K_star(x, X, theta)
55 m = len(X[0])
56 M = zeros((m,m))
57 for i in range(m):
58     M[i][i] = 1/(theta[i+2]*theta[i+2])
59 K_starstar = cov_func(x, x, M, theta[0])
60 TS = calc_E_TS(inv_C, K_star, y)
61 TS_conf = K_starstar - dot(K_star, dot(inv_C, K_star.transpose()))
62 print "Predicted TS: %d, Confidence in prediction: %d" % (TS, TS_conf)
63 return TS, TS_conf
64
65
66
67
68 #Function is necessary as for practical reasons because the training data set cannot be allowed to grow indefinitely.
    Instead
69 #new predicted values have to be swapped in for older training examples. Before they can be considered however new
    training
70 #data points have to be cross validated against the user input log to make sure they were not over ridden. This
    function takes each
71 #entry in the user log in turn and removes any entry in potential data entry that is within a certain time the user
    entry. All potential
72 #training data points that pass this test are passed into temp to be inserted into the training data set by insert_td.
73 def validate_TD(sig_time):
74     print("Validating potential new training data against user log...")
75     db = MySQLdb.connect("127.0.0.1","mike","ChickenKorma","GPR_Training_Data")
76     cur = db.cursor()
77
78     clean_time = 86400
79
80     #Import the potential training data
81     cur.execute("SELECT * FROM potential_new_training_data")
82     rows = cur.fetchall()
83     if(len(rows) > 0):
84         ptd = zeros((len(rows), len(rows[0])))
85         for j in range(len(rows)):
86             for i in range(len(rows[0])):
87                 ptd[j][i] = rows[j][i]
88
89     #Import the user log
90     cur.execute("SELECT * FROM user_TS_input_log")
91     rows = cur.fetchall()
92     user_log = zeros((len(rows), len(rows[0])))
93     for j in range(len(rows)):
94         for i in range(len(rows[0])):
95             user_log[j][i] = rows[j][i]
96
97     #PROCESS
98     for j in range(len(ptd)):
99         val_pass = True
100         for i in range(len(user_log)):
101             time_diff = (user_log[i][0] - ptd[j][6])
102             print time_diff
103             if(time_diff < 0):
104                 val_pass = True
105             elif((time_diff > 0) & (time_diff<(sig_time*60))):
106                 val_pass = False
107                 break
108             elif(time_diff > (sig_time*60)):
109                 val_pass = True
110         print("Validation test result: ", val_pass)
111
112         if(val_pass == True):
113             print("Entry in potential new training data with following id passed: ", j)
114             command = "INSERT INTO temp (thermostat_setting, external_temp, external_humidity, internal_humidity,
wind_speed, wind_direction) VALUES ('%s', '%s', '%s', '%s', '%s', '%s')%(ptd[j][0], ptd[j][1], ptd[j][2], ptd[j]
][3], ptd[j][4], ptd[j][5])"
115             print("Sending following command to database to add entry to table 'temp': ", command)
116             cur.execute(command)
117
118     print("Operation done, all entries that passed validation test moved to 'temp' for insertion into 'training data'.
")
119     cur.execute("TRUNCATE TABLE potential_new_training_data")
120     db.close()
121
122     print("Removing old data entries in user log")
123     #Get current timestamp (ms)
124     current_timestamp = int(time.time())
125     time_bound = current_timestamp - clean_time
126     command = "DELETE FROM user_TS_input_log WHERE time_stamp < '%d'" % int(time_bound)
127     cur.execute(command)
128     print("Validation process complete")
129
130
131 #Function is necessary as for practical reasons because the training data set cannot be allowed to grow indefinitely.
    Instead
132 #new predicted values have to be swapped in for older training examples. This function takes all the data points that

```

```

133     pass
134 #the validation test and adds as many of them as it can to the training set till capacity is reached. After this new
    training examples
135 #are swapped in by removing the training points that are closest in the feature space.
136 def insert_new_td(desired_training_set_size):
137     print("INSERTING NEW DATA INTO THE TRAINING DATA SET...")
138     db = MySQLdb.connect("127.0.0.1","mike","ChickenKorma","GPR_Training_Data")
139     cur = db.cursor()
140
141     #Import the number of new entries
142     cur.execute("SELECT COUNT(*) FROM temp")
143     result=cur.fetchone()
144     numb_new_entries = result[0]
145     print("Number of new entries: ", numb_new_entries)
146
147     #Import new data to be added to training data if it exists
148     if(numb_new_entries > 0):
149         cur.execute("SELECT * FROM temp")
150         rows = cur.fetchall()
151         temp = zeros((len(rows), len(rows[0])))
152         for j in range(len(rows)):
153             for i in range(len(rows[0])):
154                 temp[j][i] = rows[j][i]
155         print("Data to be inserted: ", temp)
156
157     #Import the current training data
158     cur.execute("SELECT * FROM training_data")
159     rows = cur.fetchall()
160     numrows = len(rows)
161     numcols = len(rows[0])
162     X = zeros((numrows, numcols))
163     for j in range((numrows)):
164         for i in range(numcols):
165             X[j][i] = rows[j][i]
166     X_norm = normalize_data(X)#Normalize to ensure each feature has equal significance in the sum
167     numb_training_points = len(X)
168
169     if(numb_new_entries <= abs(desired_training_set_size - numb_training_points)):#Automatically insert all points if
        there is enough capacity in the training set.
170         print("TRAINING SET NOT AT CAPACITY - INSERTING ALL NEW POINTS")
171         for j in range(len(temp)):
172             cur.execute("INSERT INTO training_data (thermostat_setting, external_temp, external_humidity,
                internal_humidity, wind_speed, wind_direction) VALUES (%s, %s, %s, %s, %s, %s)", (temp[j][0], temp[j][1], temp[j
                ][2], temp[j][3], temp[j][4], temp[j][5]))
173
174     else:
175         print("TRAINING SET NEAR CAPACITY - REPLACING OBSOLETE POINTS")#Automatically insert new points till capacity is
            reached
176         numb_auto_insert = abs(desired_training_set_size - numb_training_points)
177         print("Number to auto-insert into the training data set: ", numb_auto_insert)
178         if(numb_auto_insert > 0):
179             for j in range(numb_auto_insert):
180                 cur.execute("INSERT INTO training_data (thermostat_setting, external_temp, external_humidity,
                    internal_humidity, wind_speed, wind_direction) VALUES (%s, %s, %s, %s, %s, %s)", (temp[j][0], temp[j][1], temp[j
                    ][2], temp[j][3], temp[j][4], temp[j][5]))
181
182             for j in range(numb_auto_insert, len(temp)):#Swap in new training points for old ones that are closest to them
                in the feature space.
183                 similarities = zeros(len(X))
184                 for i in range(len(X_norm)):
185                     similarities[i] = calc_similarity(temp[j], X_norm[i])
186                 index_replace = similarities.argmax()
187                 command = "UPDATE 'training_data' SET 'thermostat_setting'='%s', 'external_temp'='%s', 'external_humidity'='%s'
                    , 'internal_humidity'='%s', 'wind_speed'='%s', 'wind_direction'='%s' WHERE 'thermostat_setting'='%s' AND '
                    external_temp'='%s' AND 'external_humidity'='%s' AND 'internal_humidity'='%s' AND 'wind_speed'='%s' AND '
                    wind_direction'='%s';" % (temp[j][0], temp[j][1], temp[j][2], temp[j][3], temp[j][4], temp[j][5], X[index_replace
                    ][0], X[index_replace][1], X[index_replace][2], X[index_replace][3], X[index_replace][4], X[index_replace][5])
188                 cur.execute(command)
189                 print("Sent following command to sequel database: ", command)
190             cur.execute("TRUNCATE TABLE temp")
191         db.close()
192
193
194 #GPR TRAIN:
195 #Function trains the GPR by optimizing hyperparameters using gradient descent.
196 def train_GPR(alpha, max_iters, error_level, init_lower, init_upper, num_init):
197     print("TRAINING GPR ALGORITHM")
198     #READ IN TRAINING DATA
199     # Open database connection
200     db = MySQLdb.connect("127.0.0.1","mike","ChickenKorma","GPR_Training_Data" )
201     # prepare a cursor object using cursor() method
202     cur = db.cursor()
203     #Read in training data from SQL database
204     cur.execute("SELECT * FROM training_data")
205     rows = cur.fetchall()
206     # disconnect from server
207     db.close()
208
209     #PUT DATA INTO MATRICES
210     #Find the number of rows (data points) and columns (Features)
211     numrows = len(rows)
212     numcols = len(rows[0])

```

```

213 #Process data into input feature data matrix X and actual outputs associated y
214 X = zeros((numrows, (numcols-1)))
215 y = zeros((numrows, 1))
216 for j in range((numrows)):
217     y[j] = rows[j][0]
218
219 for j in range(numrows):
220     for i in range(1, numcols):
221         X[j][(i-1)] = rows[j][i]
222
223 #NORMALIZE DATA
224 X = normalize_data(X)
225
226 #OPTIMISE AND DEFINE HYPERPARAMETERS USING GRADIENT DESCENT
227 #Number of Input features
228 m = len(X[0])
229 #Number of data points
230 n = len(y)
231
232 theta_init = random.uniform(init_lower,init_upper,(m+2))
233 theta, successful, nlml, nlml_history, exit_iters = gradient_descent(alpha, theta_init, X, y, error_level, max_iters
    )
234
235 #Carry out gradient descenet multiple times as with different initialisations may get different local minima and
    therefore
236 #worse or better performance. Select the set of hyperparameters that gives the lowest NLML.
237 for j in range(num_init):
238     print("j is: ", j)
239     #Define and initialise theta, note is a row vector
240     theta_init = random.uniform(init_lower,init_upper,(m+2))
241     print("Initial Theta values: ", theta_init)
242     #Use gradient descent to find optimized parameters
243     theta_test, successful, nlml_test, nlml_history, exit_iters = gradient_descent(alpha, theta_init, X, y,
        error_level, max_iters)
244     print("Initial nlml:", nlml_history[0])
245     print("Final nlml: ", nlml)
246     if(nlml_test < nlml):
247         theta = theta_test
248         nlml = nlml_test
249 print("Theta optimized: ", theta)
250
251
252 #CALCULATE MODEL VALUES
253 #Calculate K with optimized hyperparameters
254 K = calc_K(X, theta)
255 #Calculate 'C'
256 C = calc_C(K, theta[1])
257 #Calculate inverse of C
258 inv_C = linalg.inv(C)
259 #OUTPUT INVERSE C AND THETA VALUES TO DATABASE TABLES
260 # Open database connection
261 db = MySQLdb.connect("127.0.0.1","mike","ChickenKorma","GPR_Training_Data" )
262 cursor = db.cursor()
263 #Empty table of obsolete theta values
264 cursor.execute("TRUNCATE TABLE trained_theta_values")
265 #Insert trained theta values into the database
266 cursor.execute("INSERT INTO trained_theta_values (signal_power, noise_power, external_temp, external_humidity,
    internal_humidity, wind_speed, wind_direction) VALUES (%s, %s, %s, %s, %s, %s, %s)", (theta[0], theta[1], theta
    [2], theta[3], theta[4], theta[5], theta[6]))
267 #Empty table of obsolete inv_C values
268 cursor.execute("TRUNCATE TABLE inv_C")
269 #Insert inv_C values into the database
270 for j in range(n):
271     for i in range(n):
272         cursor.execute("INSERT INTO inv_C (row_number, column_number, value) VALUES (%s, %s, %s)", (j, i, inv_C[j][i]))
273 db.close()
274 print("GPR Algorithm trained")
275
276
277 #Function is called everytime a thermostat setting is requested.
278 #Description: Function reads in input data and recent user input tables. If the predicted state is occupied, and there
    has been a user input (i.e. a user input since the last call of this function) then
279 #the user input is used, otherwise former inputs from the user (found in the user log table) are used if they are
    within a certain time range of the current timestamp. Else a TS value is generated by
280 #the GPR algorithm. If the state is 'empty' or 'asleep' default settings are adopted.
281 def get_TS(internal_humidity):
282     #DEFINE VARIABLES USED IN PROGRAM
283     #k is time in minutes that a user update holds for
284     alpha = 0.0000001
285     error_level = 1
286     max_iters = 100
287     numb_init = 10
288     k = 60
289     unoccupied_temp = 2
290     desired_training_set_size = 500
291     #Initialise TS
292     TS = 0
293
294     #GET HYPERPARAMETER INITIALISATION RANGE
295     db = MySQLdb.connect("127.0.0.1","mike","ChickenKorma","algorithm_testing" )
296     cur = db.cursor()
297     cur.execute("SELECT * FROM hyperparameter_range")
298     row = cur.fetchall()

```

```

299 lb = row[0][0]
300 ub = row[0][1]
301 db.close()
302
303 #GET DATA REQUIRED
304 db = MySQLdb.connect("127.0.0.1","mike","ChickenKorma","Data" )
305 cur = db.cursor()
306 #Import user settings
307 cur.execute("SELECT * FROM limits")
308 row = cur.fetchall()
309 max_temp = row[0][0]
310 min_temp = row[0][1]
311 sleeping_temp = row[0][2]
312 print("IMPORTED USER SETTINGS: ")
313 print("max_temp is: ", max_temp)
314 print("min_temp is: ", min_temp)
315 print("sleeping_temp is: ", sleeping_temp)
316
317 #Import the input for this call
318 x, predicted_state, current_timestamp = get_input(internal_humidity)
319
320 print("DATA TO BE USED TO GENERATE TS IF NO USER INPUT")
321 print("New input data entry: ",x)
322 print("State: ",predicted_state)
323 print("Timestamp: ", current_timestamp)
324
325 print("ASCERTAINING IF THERE HAS BEEN A USER UPDATE...")
326 #Ascertain if there has been a recent user update
327 cur.execute("SELECT * FROM user_update")
328 rows = cur.fetchall()
329 numb_recent_updates = len(rows)
330 print("Number of user updates in last k minutes: ", numb_recent_updates)
331
332 if(numb_recent_updates > 0):
333     new_user_update = True
334     user_update = zeros((len(rows), len(rows[0])))
335     for j in range(len(rows)):
336         for i in range(len(rows[0])):
337             user_update[j][i] = rows[j][i]
338         print("user update data: ", user_update)
339
340 else:
341     new_user_update = False
342
343
344 print("State of user update: ", new_user_update)
345 db.close()
346
347 #PROCESS
348 #If the user has made a change since the last call then use this
349 if(new_user_update == True):
350     print("USER UPDATE DETECTED - USER INPUT PATH TAKEN")
351     if(predicted_state == 0):
352         print("State: Unoccupied, setting TS to unoccupied setting...")
353         TS = unoccupied_temp
354     elif(predicted_state == 1):
355         print("State: Occupied, setting TS to user setting...")
356         TS = user_update[(numb_recent_updates-1)][0]
357     elif(predicted_state == 2):
358         print("State: Asleep, setting TS to sleeping setting...")
359         TS = sleeping_temp
360     else:
361         print("Error: 'predicted_state' value not valid:", predicted_state)
362
363
364 if(TS > max_temp):
365     TS = max_temp
366
367 elif(TS < min_temp):
368     TS = min_temp
369
370 else:
371     print("TS within boundry conditions")
372
373 print("TS value selected",TS)
374
375 #OUTPUT TS TO THE EMBED
376 print("OUTPUT TO EMBED...")
377 ws = create_connection("ws://ec2-54-214-164-65.us-west-2.compute.amazonaws.com:8888/ws")
378 print("Sending Thermostat Setting...")
379 command = '{"type": "thermostat_control" , "setting" : %d}' % int(TS)
380 print("Output to embed: ", command)
381 ws.send(command)
382 print("Thermostat Setting Sent with value of: ", TS)
383 ws.close()
384
385
386 #HANDLE DATABASE
387 db = MySQLdb.connect("127.0.0.1","mike","ChickenKorma","GPR_Training_Data" )
388 cur = db.cursor()
389 #Upload the user input to the temp file and retrain GPR algorithm
390 cur.execute("INSERT INTO temp (thermostat_setting, external_temp, external_humidity, internal_humidity, wind_speed
, wind_direction) VALUES (%, %, %, %, %, %s)", (user_update[(numb_recent_updates-1)][0], user_update[(

```



```

numb_recent_updates-1]][1], user_update[(numb_recent_updates-1)][2], user_update[(numb_recent_updates-1)][3],
user_update[(numb_recent_updates-1)][4], user_update[(numb_recent_updates-1)][5]))
391 #Upload the user input to the user log
392 command = "INSERT INTO user_TS_input_log (thermostat_setting, external_temp, external_humidity, internal_humidity,
wind_speed, wind_direction, time_stamp) VALUES (%s, %s, %s, %s, %s, %s, %s)" % (user_update[(numb_recent_updates
-1)][0], user_update[(numb_recent_updates-1)][1], user_update[(numb_recent_updates-1)][2], user_update[(
numb_recent_updates-1)][3], user_update[(numb_recent_updates-1)][4], user_update[(numb_recent_updates-1)][5],
user_update[(numb_recent_updates-1)][6])
393 print("Uploading following entry into user log: ", command)
394 cur.execute(command)
395 #Retrain Algorithm and truncate 'temp' afterwards
396 insert_new_td(desired_training_set_size)
397 train_GPR(alpha, max_iters, error_level, lb, ub, numb_init)
398 db.close()
399
400 db = MySQLdb.connect("127.0.0.1", "mike", "ChickenKorma", "Data" )
401 cur = db.cursor()
402 #Truncate user_data (user input)
403 cur.execute("TRUNCATE TABLE user_update")
404 #Truncate Data (input)
405 cur.execute("TRUNCATE TABLE Data")
406
407
408 #NO INPUT SINCE LAST CALL...
409 else:
410     print("USER UPDATE NOT DETECTED - CHECKING FOR FORMER STILL RELEVANT USER INPUTS...")
411
412     #Import user log data
413     db = MySQLdb.connect("127.0.0.1", "mike", "ChickenKorma", "GPR_Training_Data" )
414     cur = db.cursor()
415     cur.execute("SELECT * FROM user_TS_input_log")
416     rows = cur.fetchall()
417     numb_user_entries = len(rows)
418     user_time_stamps = zeros(numb_user_entries)
419
420     #Load in the time stamps in the user log
421     for j in range(numb_user_entries):
422         user_time_stamps[j] = rows[j][0]
423
424     print("Numer of entries in user log: ", numb_user_entries)
425     print("User Log timestamps:", user_time_stamps)
426
427     if(numb_user_entries > 0):
428         #Check if any entries in the user log are relevant
429         id, closest_user_timestamp, min_diff = find_nearest(user_time_stamps, current_timestamp)
430         if(fabs(current_timestamp - closest_user_timestamp) < (k*60)): #Note timestamp is in milliseconds, k*60*10^3 ms
in k minutes
431             use_user_input = True
432
433         else:
434             use_user_input = False
435
436         print("Current Time Stamp: ", current_timestamp)
437         print("Closest time difference to current timestamp: ", min_diff)
438
439     else:
440         use_user_input = False
441
442     print("State of use user log input: ", use_user_input)
443
444     #Process
445     if(predicted_state == 0):
446         print("State: Unoccupied, setting TS to unoccupied setting")
447         TS = unoccupied_temp
448     elif(predicted_state == 1):
449         if(use_user_input == True):
450             print("Former user input still relevant, referring back to last user log entry...")
451             TS = user_TS(current_timestamp)
452         else:
453             print("Former user entries no longer relevant, generating TS from GPR algorithm...")
454             TS, TS_conf = active_TS(x)
455     elif(predicted_state == 2):
456         print("State: Asleep, setting TS to sleeping setting...")
457         TS = sleeping_temp
458     else:
459         print("Error: 'predicted_state' value not valid:")
460         print(predicted_state)
461
462     #Bound the TS setting
463     if(TS > max_temp):
464         TS = max_temp
465
466     elif(TS < min_temp):
467         TS = min_temp
468     else:
469         print("TS within boundry conditions")
470
471     #OUTPUT TS TO THE EMBED
472     print("OUTPUT TO EMBED...")
473     ws = create_connection("ws://ec2-54-214-164-65.us-west-2.compute.amazonaws.com:8888/ws")
474     print("Sending Thermostat Setting...")
475     command = '{"type": "thermostat_control" , "setting" : %d}' % int(TS)
476     print("Output to embed: ", command)

```

```

477 ws.send(command)
478 print("Thermostat Setting Sent with value of: ", TS)
479 ws.close()
480
481
482 #HANDLE DATABASE UPDATES
483 if((use_user_input == False) & (predicted_state == 1)):
484     db = MySQLdb.connect("127.0.0.1","mike","ChickenKorma","GPR_Training_Data" )
485     cur = db.cursor()
486     #Upload input to potential new training data
487     cur.execute("INSERT INTO potential_new_training_data (thermostat_setting, external_temp, external_humidity,
internal_humidity, wind_speed, wind_direction, time_stamp) VALUES (%s, %s, %s, %s, %s, %s, %s)", (TS, x[0][0], x
[0][1], x[0][2], x[0][3], x[0][4], current_timestamp))
488     db.close()
489
490
491
492
493 #Function takes all the data in the potential_training_data table, runs a validation check against the
user_TS_input_log and then outputs the data entries
494 #that pass into the 'temp' data table these are then inserted into the training data table and the GPR retrain
algorithm is re-run. The temp and potential data table are then truncated
495 #and all entries older than a day in the user log are also removed. Note user inputs are automatically added to the
training set as soon as they are detected. Called approximatly
496 #every 24 hours.
497 def retrain_with_updates():
498     #GET HYPERPARAMETER INITIALISATION RANGE
499     db = MySQLdb.connect("127.0.0.1","mike","ChickenKorma","algorithm_testing" )
500     cur = db.cursor()
501     cur.execute("SELECT * FROM hyperparameter_range")
502     row = cur.fetchall()
503     lb = row[0][0]
504     ub = row[0][1]
505     db.close()
506
507     #PARAMETER DEFINITIONS
508     alpha = 0.0000001
509     error_level = 1
510     max_iters = 100
511     numb_init = 10
512     sig_time = 60
513     desired_training_set_size = 500
514
515     #Check data against user log entries, if within a certain time distance then don't insert into the table 'temp'.
Once complete truncates 'potential_training_data' and removes old entries in user log
516     validate_TD(sig_time)
517
518     #Insert data in 'temp' table into the 'training_data' table appropriately. Once complete truncates 'temp'
519     insert_new_td(desired_training_set_size)
520
521     #Retrain Algorithm, i.e. obtain new theta values and correlation matrix
522     train_GPR(alpha, max_iters, error_level, lb, ub, numb_init)
523
524
525
526 #HYPERPARAMETER SEARCH:
527 #Different initialisations of theta lead to different local minima and therefore different performance. This function
searches the NLML space and initializes
528 #the hyperparameters in sub regions. It checks the performance of each region by comparing the predictions with the
actual values from a test set (used to train
529 #the data) and a cross validation set (an independant set) and calculates the percentage error for both. The region
with the best performance is then subdivided again
530 #until a certain initialisation range is reached. It then returns this initialisation range to the database to be used
by all future 'train calls' to initialise the hyperparameters.#
531 #This function is simply put a range finder and gives an idea of where to look for the best minima, it is only called
when there are very large changes in the data set.
532 def hyperparameter_search():
533     #INITIALISE SYSTEM PARAMETERS
534     alpha = 0.00001
535     error_level = 1
536     max_iters = 50
537     num_init = 20
538
539     #IMPORT INPUT DATA
540     db = MySQLdb.connect("127.0.0.1","mike","ChickenKorma","algorithm_testing" )
541     cur = db.cursor()
542
543     #Read in training data
544     cur.execute("SELECT * FROM training_data")
545     rows = cur.fetchall()
546     train_d = zeros((len(rows), len(rows[0])))
547     for j in range(len(rows)):
548         for i in range(len(rows[0])):
549             train_d[j][i] = rows[j][i]
550
551     #Read in test data
552     cur.execute("SELECT * FROM test_data")
553     rows = cur.fetchall()
554     test_d = zeros((len(rows), (len(rows[0])-1)))
555     TS_test = zeros((len(rows)))
556     #print len(rows)
557     for j in range(len(rows)):
558         TS_test[j] = rows[j][0]

```

```

559     for i in range(1,len(rows[0])):
560         test_d[j][(i-1)] = rows[j][i]
561
562     lower_bound = 0
563     upper_bound = 100
564     numb_divs = 10
565     m = len(test_d[0])
566
567     #Perform the grid search
568     while((upper_bound - lower_bound) > 4):
569         increment = (upper_bound - lower_bound)/numb_divs
570         print "Considering range of initialisations %d to %d:" % (lower_bound, upper_bound)
571         print "Increment is: %d" % (increment)
572         lb_temp = lower_bound
573         ub_temp = lb_temp + increment
574         av_err = zeros((numb_divs))
575         for k in range((numb_divs)):
576             #TRAIN ALGORITHM:
577             train_GPR(alpha, max_iters, error_level, lb_temp, ub_temp, num_init)
578             #Initialise
579             err = zeros((len(TS_test)))
580             gpr_res = zeros((len(TS_test)))
581             gpr_conf = zeros((len(TS_test)))
582             for j in range(len(test_d)):
583                 gpr_res[j], gpr_conf[j] = active_TS(test_d[j])
584                 err[j] = 100*((abs(TS_test[j] - gpr_res[j]))/TS_test[j])
585                 av_err[k] = av_err[k] + err[j]
586
587             av_err[k] = av_err[k]/(len(test_d))
588             print "Error bound for division %d-%d is: %d" % (lb_temp, ub_temp, av_err[k])
589             lb_temp = ub_temp
590             ub_temp = ub_temp + increment
591
592             ind_min_err = np.argmin(av_err)
593             print "the index of the minimum division is: %d" % (ind_min_err)
594             upper_bound = lower_bound + (ind_min_err*increment)
595             lower_bound = upper_bound - increment
596             print "Theta values found to be best initialised in this range are between %d - %d" % (lower_bound, upper_bound)
597
598
599     print "Theta values found to be best initialised in this range are between %d - %d" % (lower_bound, upper_bound)
600     print("Initialised with these Boundaries we have...")
601
602     train_GPR(alpha, max_iters, error_level, lower_bound, upper_bound, num_init)
603     err = zeros((len(TS_test)))
604     gpr_res = zeros((len(TS_test)))
605     gpr_conf = zeros((len(TS_test)))
606     av_err = 0
607     for j in range(len(test_d)):
608         gpr_res[j], gpr_conf[j] = active_TS(test_d[j])
609         err[j] = 100*((abs(TS_test[j] - gpr_res[j]))/TS_test[j])#Calculate error on each prediction
610         av_err = av_err + err[j]
611     av_err = av_err/len(test_d)#Calculate the % error in the test set
612
613     min = np.argmin(err)
614     max = np.argmax(err)
615
616     print("TS values from test set", TS_test)
617     print("TS values calculated", gpr_res)
618     print("Variance of GPR values", gpr_conf)
619     print("Percentage Error List:", err)
620     print "Max percentage error in estimate: %d" % (err[max])
621     print "Min percentage error in estimate: %d" % (err[min])
622     print "Average Percentage Error in estimates: %d" % (av_err)
623
624     cur.execute("TRUNCATE TABLE hyperparameter_range")
625     cur.execute("INSERT INTO hyperparameter_range (lower_bound, upper_bound) VALUES (%s, %s)" % (lower_bound,
        upper_bound))
626
627
628
629
630
631
632     #With insufficient data to run the state detection algorithm this function is used to generate values
633     def estimate_current_state():
634         x = random.randint(1,100)
635         if x <= 10:
636             return 0
637         elif x <= 66:
638             return 1
639         else:
640             return 2
641
642
643
644     #Function is run before 'get_TS' is called to compile an input vector
645     def get_input(internal_humidity):
646         db = MySQLdb.connect("127.0.0.1","mike","ChickenKorma","Data" )
647         cur = db.cursor()
648         #Import enviromental data from Data2
649         command = "SELECT Temperature, Humidity, Wind_speed, Wind_direction FROM Data2 WHERE Time_stamp=(select MAX(
            Time_stamp) from Data2)"

```

```

650 print("Command sent to SQL database: ", command)
651 cur.execute(command)
652 temp = cur.fetchone()
653 print temp
654 #Get the state of the house:
655 state = find_state()
656 #Get current Time stamp
657 current_timestamp = int(time.time())/1000
658 #Format Oututs:
659 x = zeros((5))
660 x[0] = temp[0]
661 x[1] = temp[1]
662 x[2] = internal_humidity
663 x[3] = temp[2]
664 x[4] = temp[3]
665 return x, state, current_timestamp
666
667
668
669
670 #Function uses a logic table to select a state by comparing the predicted and detected state. It then returns this
    state to
671 #the 'get_TS' function.
672 def find_state():
673     localtime = time.localtime(time.time())
674     print "Day is: %d" % (localtime[6])
675     print "Hour is: %d" % (localtime[3])
676
677     hour = localtime[3]
678
679     db = MySQLdb.connect("127.0.0.1","mike","ChickenKorma","House_states" )
680     cur = db.cursor()
681
682     #First need find the correct day, the data is puleld of the same table that feeds the website.
683     if(localtime[6] == 0):
684         command = "SELECT state FROM Monday WHERE Hour = '%d'" % (hour)
685         print("Command sent to SQL database: ", command)
686         cur.execute(command)
687         temp = cur.fetchone()
688         pred_state = int(temp[0])
689
690     elif(localtime[6] == 1):
691         command = "SELECT state FROM Tuesday WHERE Hour = '%d'" % (hour)
692         print("Command sent to SQL database: ", command)
693         cur.execute(command)
694         temp = cur.fetchone()
695         pred_state = int(temp[0])
696
697     elif(localtime[6] == 2):
698         command = "SELECT state FROM Wednesday WHERE Hour = '%d'" % (hour)
699         print("Command sent to SQL database: ", command)
700         cur.execute(command)
701         temp = cur.fetchone()
702         pred_state = temp[0]
703
704     elif(localtime[6] == 3):
705         command = "SELECT state FROM Thursday WHERE Hour = '%d'" % (hour)
706         print("Command sent to SQL database: ", command)
707         cur.execute(command)
708         temp = cur.fetchone()
709         pred_state = temp[0]
710
711     elif(localtime[6] == 4):
712         command = "SELECT state FROM Friday WHERE Hour = '%d'" % (hour)
713         print("Command sent to SQL database: ", command)
714         cur.execute(command)
715         temp = cur.fetchone()
716         pred_state = temp[0]
717
718     elif(localtime[6] == 5):
719         command = "SELECT state FROM Saturday WHERE Hour = '%d'" % (hour)
720         print("Command sent to SQL database: ", command)
721         cur.execute(command)
722         temp = cur.fetchone()
723         pred_state = temp[0]
724
725     elif(localtime[6] == 6):
726         command = "SELECT state FROM Sunday WHERE Hour = '%d'" % (hour)
727         print("Command sent to SQL database: ", command)
728         cur.execute(command)
729         temp = cur.fetchone()
730         pred_state = temp[0]
731     else:
732         print "Day not recognized, day = '%d'" % (localtime[6])
733     db.close()
734
735     rt_state = estimate_current_state()
736     print type(rt_state)
737     print type(pred_state)
738 #Logic table
739 if(rt_state == 0 and (pred_state == 0 or pred_state == 3)):
740     state = 0
741 elif(rt_state == 0 and (pred_state == 1 or pred_state == 4)):

```

```

742     state = 1
743     elif(rt_state == 0 and (pred_state == 2 or pred_state == 5)):
744         state = 0
745     elif(rt_state == 1 and (pred_state == 0 or pred_state == 3)):
746         state = 1
747     elif(rt_state == 1 and (pred_state == 1 or pred_state == 4)):
748         state = 1
749     elif(rt_state == 1 and (pred_state == 2 or pred_state == 5)):
750         state = 2
751     elif(rt_state == 2 and (pred_state == 0 or pred_state == 3)):
752         state = 2
753     elif(rt_state == 2 and (pred_state == 1 or pred_state == 4)):
754         state = 1
755     elif(rt_state == 2 and (pred_state == 2 or pred_state == 5)):
756         state = 2
757     else:
758         print "Real time and predicted state not recognised, rt_state = '%d', pred_state = '%d'" % (rt_state, pred_state)
759
760     return state
761
762 random.seed()

```

code/gpr_system_primary_functions.py

8.3 Costs

Component	Price/unit	Quantity	Total price
Wifly module	£22.41	1	£22.41
Mbed	£41.38	1	£41.38
Motion sensor	£12.24	1	£12.24
Zigbee	£17.55	4	£70.20
Temperature Sensor	£1.21	1	£1.21
Breadboard	£3.15	5	£15.75
XBee to DIP Adapter	£2.50	4	£10.00
Humidity Sensor	£3.42	1	£3.42
Amega88 MCU	£2.66	1	£2.66
9V AC-AC adapter	£6.25	1	£6.25
2.1mm socket	£0.65	1	£0.65
Triac	£2.70	2	£5.40
Optocoupler	£1.07	1	£1.07
TL072 opamp	£0.68	2	£1.36
100W SPOT ES	£0.32	1	£0.32
Comparator	£0.17	4	£0.68
XOR gate	£0.26	2	£0.52
DC power adapter	£8.52	1	£8.52
Lead kettle	£2.20	1	£2.20
Voltage regulator	£0.60	1	£0.60
Small breadboard	£3.00	1	£3.00
Voltage regulator	£0.60	2	£1.20
2.1mm socket	£0.65	2	£1.30
		Total	£212.34